

# Adult Census Income

June 24, 2019

## 1 Adult Census Income

Predict whether income exceeds \$50K/yr based on census data

---

### 1.1 Informations on the dataset

This data was extracted from the 1994 Census bureau database by Ronny Kohavi and Barry Becker (Data Mining and Visualization, Silicon Graphics). A set of reasonably clean records was extracted using the following conditions: ((AAGE>16) && (AGI>100) && (AFNLWGT>1) && (HRSWK>0)). The prediction task is to determine whether a person makes over \$50K a year.

Original dataset open sourced, can be found [here](#).

### 1.2 Goal

Predict **whether or not a person makes more than USD 50,000** from the information contained in the columns. Find clear insights on the profiles of the people that make more than 50,000USD / year. For example, which variables seem to be the most correlated with this phenomenon?

---

## 2 Dataset first insight

Libraries import

```
In [1]: import warnings
        warnings.simplefilter(action='ignore', category=FutureWarning)

In [2]: import numpy as np
        import pandas as pd
        import matplotlib.pyplot as plt
        import seaborn as sns

In [3]: from sklearn.model_selection import train_test_split
        from sklearn.metrics import accuracy_score

        from sklearn.linear_model import LogisticRegression
```

```

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, ExtraTreesClassifier
from sklearn.model_selection import GridSearchCV

```

Loading the file

```

In [4]: df = pd.read_csv('./input/adult.csv')
df.head()

```

```

Out[4]:
   age  workclass  fnlwgt  education  education.num  marital.status  \
0   90         ?   77053    HS-grad             9         Widowed
1   82   Private  132870    HS-grad             9         Widowed
2   66         ?  186061  Some-college          10         Widowed
3   54   Private  140359     7th-8th             4         Divorced
4   41   Private  264663  Some-college          10         Separated

   occupation  relationship  race  sex  capital.gain  \
0           ?  Not-in-family  White  Female           0
1  Exec-managerial  Not-in-family  White  Female           0
2           ?    Unmarried  Black  Female           0
3  Machine-op-inspct    Unmarried  White  Female           0
4   Prof-specialty    Own-child  White  Female           0

   capital.loss  hours.per.week  native.country  income
0           4356             40  United-States  <=50K
1           4356             18  United-States  <=50K
2           4356             40  United-States  <=50K
3           3900             40  United-States  <=50K
4           3900             40  United-States  <=50K

```

Columns description

- age: continuous.
- workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.
- fnlwgt: continuous.
- education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.
- education-num: continuous.
- marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.
- occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.
- relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.
- race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.
- sex: Female, Male.
- capital-gain: continuous.
- capital-loss: continuous.

- hours-per-week: continuous.
- native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.

```
In [5]: df.shape
```

```
Out[5]: (32561, 15)
```

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
age                32561 non-null int64
workclass          32561 non-null object
fnlwgt             32561 non-null int64
education          32561 non-null object
education.num      32561 non-null int64
marital.status     32561 non-null object
occupation         32561 non-null object
relationship       32561 non-null object
race               32561 non-null object
sex               32561 non-null object
capital.gain       32561 non-null int64
capital.loss       32561 non-null int64
hours.per.week     32561 non-null int64
native.country     32561 non-null object
income            32561 non-null object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB
```

When it comes to numerical values, no information is missing. On the contrary for categorical features, there are '?', which indicated unknow information. Some rows are duplicated and need to be removed :

```
In [7]: df.duplicated().sum()
```

```
Out[7]: 24
```

```
In [8]: df = df.drop_duplicates()
        df.shape
```

```
Out[8]: (32537, 15)
```

```
In [9]: cat_feat = df.select_dtypes(include=['object']).columns
cat_feat
```

```
Out[9]: Index(['workclass', 'education', 'marital.status', 'occupation',
              'relationship', 'race', 'sex', 'native.country', 'income'],
              dtype='object')
```

The number of missing value isn't relevant

```
In [10]: print('% of missing values :')
         for c in cat_feat:
             perc = len(df[df[c] == '?']) / df.shape[0] * 100
             print(c, f'{perc:.1f} %')
```

```
% of missing values :
workclass 5.6 %
education 0.0 %
marital.status 0.0 %
occupation 5.7 %
relationship 0.0 %
race 0.0 %
sex 0.0 %
native.country 1.8 %
income 0.0 %
```

Basic statistics for numerical values:

```
In [11]: df.describe()
```

```
Out[11]:
```

	age	fnlwgt	education.num	capital.gain	capital.loss	\
count	32537.000000	3.253700e+04	32537.000000	32537.000000	32537.000000	
mean	38.585549	1.897808e+05	10.081815	1078.443741	87.368227	
std	13.637984	1.055565e+05	2.571633	7387.957424	403.101833	
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	
25%	28.000000	1.178270e+05	9.000000	0.000000	0.000000	
50%	37.000000	1.783560e+05	10.000000	0.000000	0.000000	
75%	48.000000	2.369930e+05	12.000000	0.000000	0.000000	
max	90.000000	1.484705e+06	16.000000	99999.000000	4356.000000	

	hours.per.week
count	32537.000000
mean	40.440329
std	12.346889
min	1.000000
25%	40.000000
50%	40.000000
75%	45.000000
max	99.000000

### 3 Exploratory Analysis

```
In [12]: # Taking a look at the target (income) without distinction of sex
         print(f"Ratio above 50k : {(df['income'] == '>50K').astype('int').sum() / df.shape[0]}")
```

Ratio above 50k : 24.09%

Distinction between numerical vs. text values

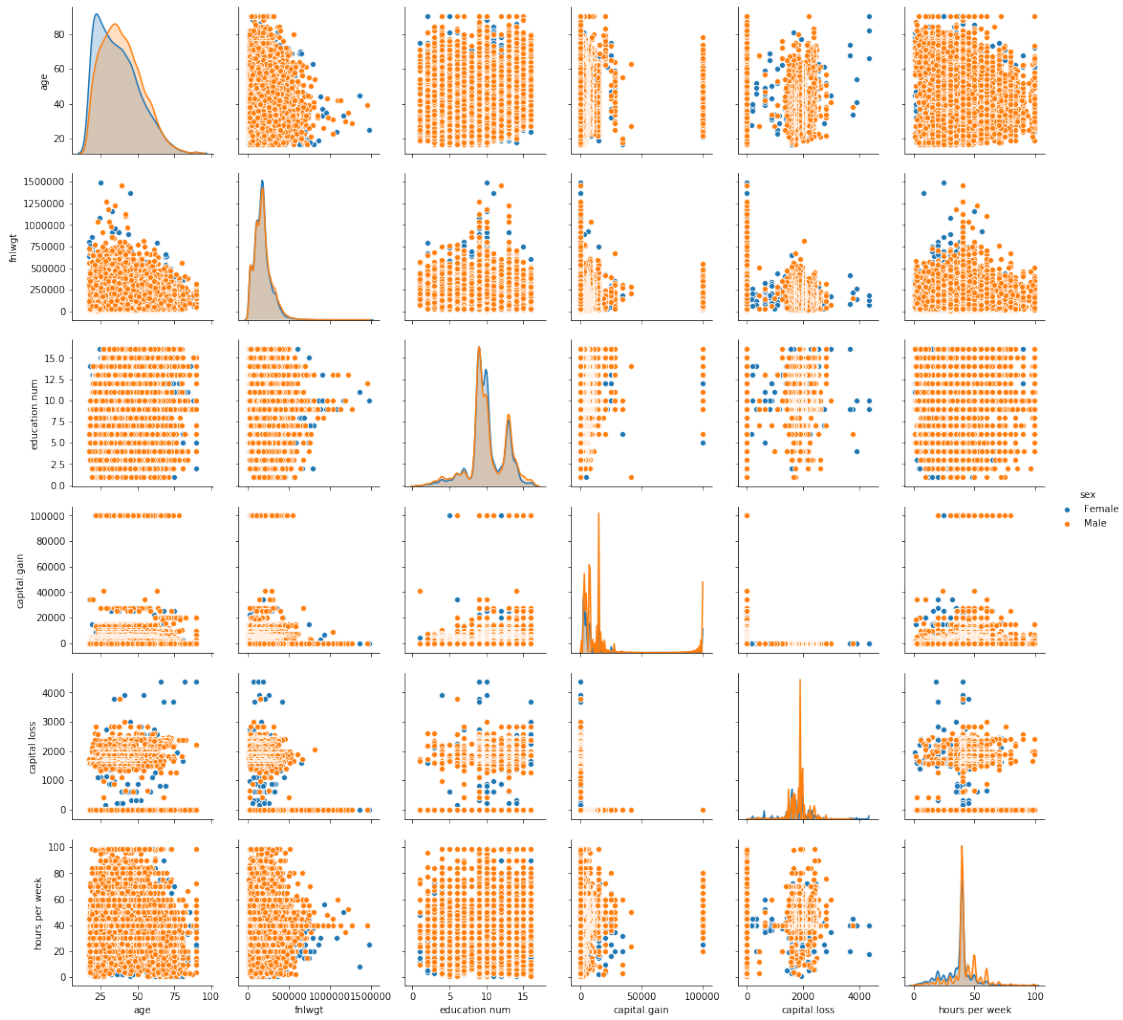
```
In [13]: num_feat = df.select_dtypes(include=['int64']).columns
         num_feat
```

```
Out[13]: Index(['age', 'fnlwgt', 'education.num', 'capital.gain', 'capital.loss',
               'hours.per.week'],
              dtype='object')
```

Plot pairwise relationships in a dataset.

```
In [14]: plt.figure(1, figsize=(16,10))
         sns.pairplot(data=df, hue='sex')
         plt.show()
```

<Figure size 1152x720 with 0 Axes>



Distributions of numerical values

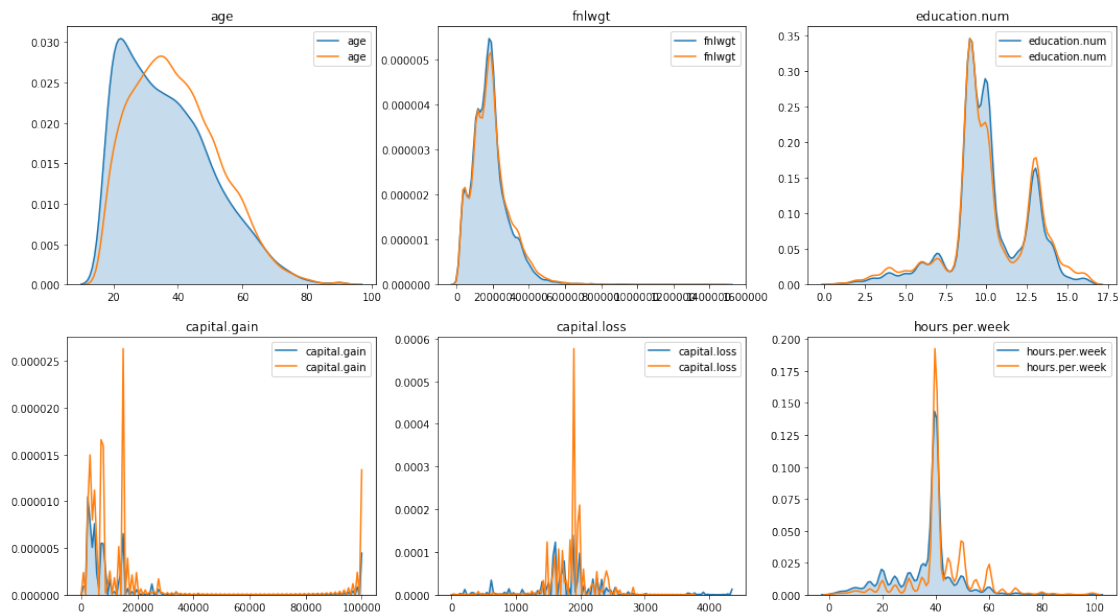
```
In [15]: plt.figure(figsize=(18,10))
plt.subplot(231)
```

```
i=0
for c in num_feat:
    plt.subplot(2, 3, i+1)
    i += 1
    sns.kdeplot(df[df['sex'] == 'Female'][c], shade=True, )
    sns.kdeplot(df[df['sex'] == 'Male'][c], shade=False)
    plt.title(c)

plt.show()
```

/home/sunflowa/anaconda3/lib/python3.7/site-packages/matplotlib/figure.py:98: MatplotlibDeprecationWarning: Adding an axes using the same arguments as a previous axes currently reuses the earlier instance

"Adding an axes using the same arguments as a previous axes "



There are significant differences when it comes to capital gain / loss and hours per week.

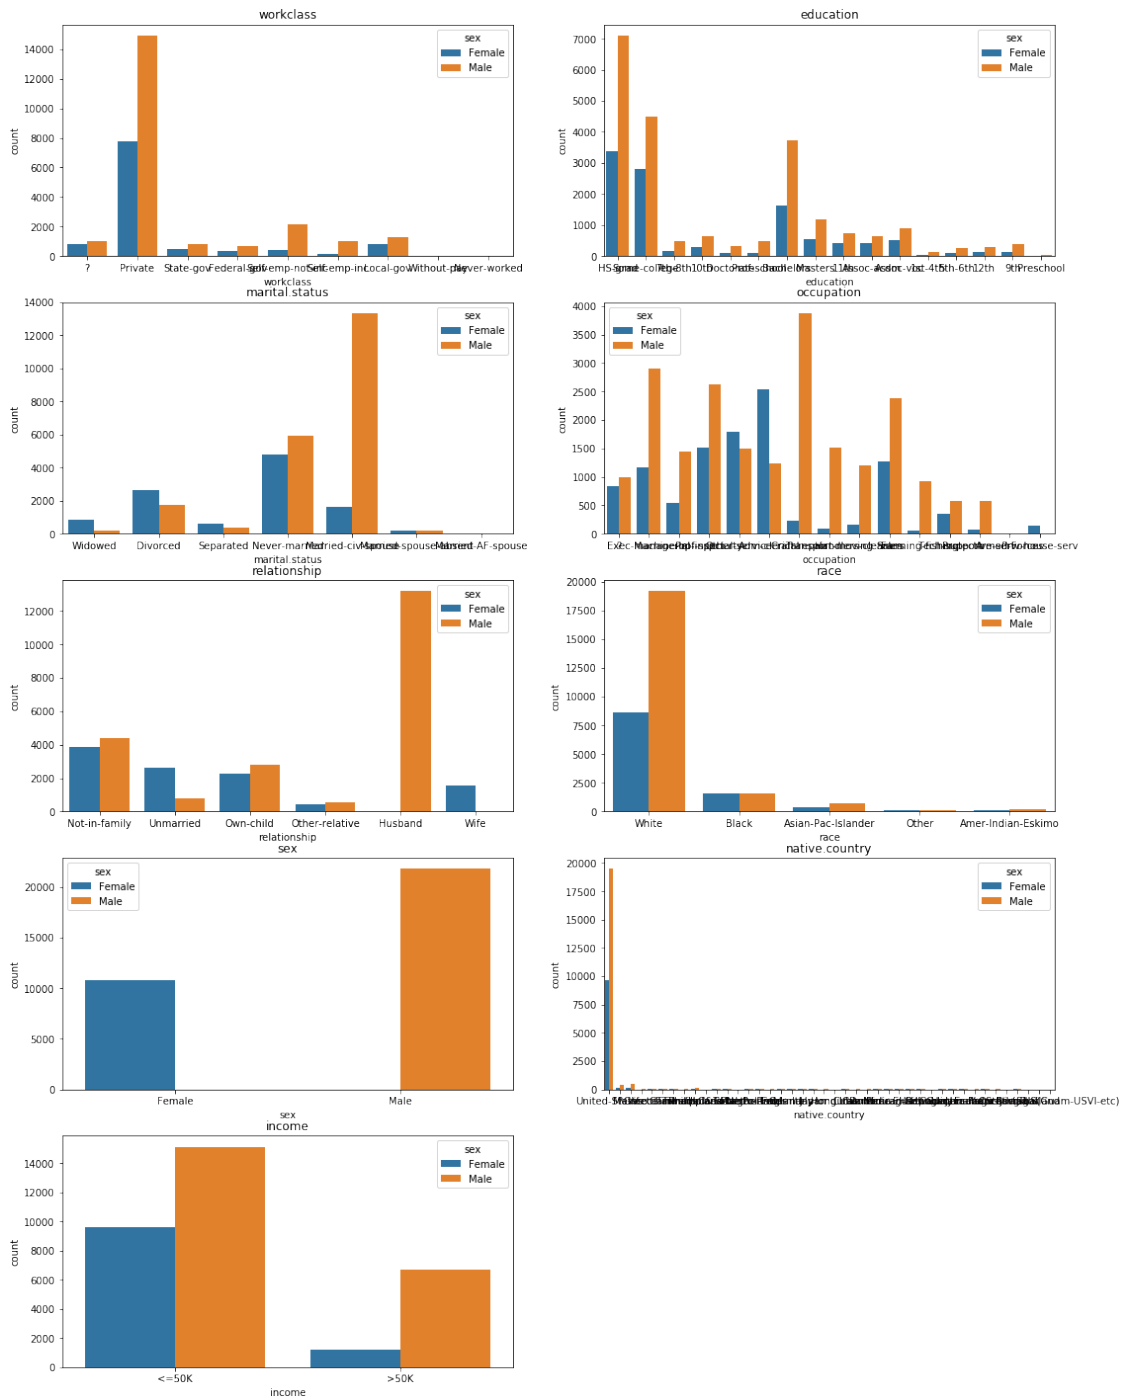
```
In [16]: plt.figure(figsize=(18,25))
plt.subplot(521)
```

```
i=0
for c in cat_feat:
    plt.subplot(5, 2, i+1)
    i += 1
    sns.countplot(x=c, data=df, hue='sex')
    plt.title(c)
```

```
plt.show()
```

/home/sunflowa/anaconda3/lib/python3.7/site-packages/matplotlib/figure.py:98: MatplotlibDeprecationWarning: Adding an axes using the same arguments as a previous axes currently reuses the earlier instance.

"Adding an axes using the same arguments as a previous axes "



There are far more male earning >50k than female, but at the same time there are also more male earning <50k and even more males recorded in general. The counts need to be normalized.

```
In [17]: # nb of female / male
nb_female = (df.sex == 'Female').astype('int').sum()
nb_male = (df.sex == 'Male').astype('int').sum()
nb_female, nb_male
```



```
Out[17]: (10762, 21775)
```

```
In [18]: # nb of people earning more or less than 50k per gender
nb_male_above = len(df[(df.income == '>50K') & (df.sex == 'Male')])
nb_male_below = len(df[(df.income == '<=50K') & (df.sex == 'Male')])
nb_female_above = len(df[(df.income == '>50K') & (df.sex == 'Female')])
nb_female_below = len(df[(df.income == '<=50K') & (df.sex == 'Female')])
nb_male_above, nb_male_below, nb_female_above, nb_female_below
```

```
Out[18]: (6660, 15115, 1179, 9583)
```

```
In [19]: print(f'Among Males    : {nb_male_above/nb_male*100:.0f}% earn >50K // {nb_male_below/nb_male*100:.0f}% earn <=50K')
print(f'Among Females    : {nb_female_above/nb_female*100:.0f}% earn >50K // {nb_female_below/nb_female*100:.0f}% earn <=50K')
```

```
Among Males    : 31% earn >50K // 69% earn <=50K
Among Females  : 11% earn >50K // 89% earn <=50K
```

```
In [20]: # normalization
nb_male_above /= nb_male
nb_male_below /= nb_male
nb_female_above /= nb_female
nb_female_below /= nb_female
nb_male_above, nb_male_below, nb_female_above, nb_female_below
```


```
Out[20]: (0.3058553386911596,
0.6941446613088404,
0.1095521278572756,
0.8904478721427244)
```

```
In [21]: print(f'Among people earning >50K : {nb_male_above / (nb_male_above + nb_female_above)*100:.0f}% are Males and {nb_female_above / (nb_male_above + nb_female_above)*100:.0f}% are Females')
print(f'Among people earning <=50K : {nb_male_below / (nb_male_below + nb_female_below)*100:.0f}% are Males and {nb_female_below / (nb_male_below + nb_female_below)*100:.0f}% are Females')
```


```
Among people earning >50K : 74% are Females and 26% are Males
Among people earning <=50K : 44% are Females and 56% are Males
```

The distinction between american natives and the others should also be made

```
In [22]: df['US native'] = (df['native.country'] == 'United-States').astype('int')
plt.figure(figsize=(6,4))
sns.countplot(x='US native', data=df, hue='sex')
plt.show()
```



```
In [23]: plt.figure(figsize=(6,4))  
sns.countplot(x='income', data=df, hue='US native')  
plt.show()
```



```

In [24]: # nb of people earning more or less than 50k per origin
nb_native_above = len(df[(df.income == '>50K') & (df['US native'] == 1)])
nb_native_below = len(df[(df.income == '<=50K') & (df['US native'] == 1)])
nb_foreign_above = len(df[(df.income == '>50K') & (df['US native'] == 0)])
nb_foreign_below = len(df[(df.income == '<=50K') & (df['US native'] == 0)])
nb_native_above, nb_native_below, nb_foreign_above, nb_foreign_below

Out[24]: (7169, 21984, 670, 2714)

In [25]: nb_native = (df['US native'] == 1).astype('int').sum()
nb_foreign = df.shape[0] - nb_native
nb_native, nb_foreign

Out[25]: (29153, 3384)

In [26]: print(f'Among natives      : {nb_native_above/nb_native*100:.0f}% earn >50K // {nb_native_below/nb_native*100:.0f}% earn <=50K')
print(f'Among foreigners : {nb_foreign_above/nb_foreign*100:.0f}% earn >50K // {nb_foreign_below/nb_foreign*100:.0f}% earn <=50K')

Among natives      : 25% earn >50K // 75% earn <=50K
Among foreigners : 20% earn >50K // 80% earn <=50K

In [27]: # normalization
nb_native_above /= nb_native
nb_native_below /= nb_native
nb_foreign_above /= nb_foreign
nb_foreign_below /= nb_foreign
nb_native_above, nb_native_below, nb_foreign_above, nb_foreign_below

Out[27]: (0.24590951188556923,
0.7540904881144308,
0.1979905437352246,
0.8020094562647754)

In [28]: print(f'Among people earning >50K : {nb_native_above / (nb_native_above + nb_foreign_above)*100:.0f}% are natives and {nb_foreign_above / (nb_native_above + nb_foreign_above)*100:.0f}% are foreigners')
print(f'Among people earning <=50K : {nb_native_below / (nb_native_below + nb_foreign_below)*100:.0f}% are natives and {nb_foreign_below / (nb_native_below + nb_foreign_below)*100:.0f}% are foreigners')

Among people earning >50K : 55% are natives and 45% are foreigners
Among people earning <=50K : 48% are natives and 52% are foreigners

In [29]: num_feat = df.select_dtypes(include=['float', 'int']).columns
num_feat

Out[29]: Index(['age', 'fnlwgt', 'education.num', 'capital.gain', 'capital.loss',
'hours.per.week', 'US native'],
dtype='object')

```

```

In [30]: sns.set(style="white")

# Compute the correlation matrix
corr = df[num_feat].corr()

# Generate a mask for the upper triangle
mask = np.zeros_like(corr, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True

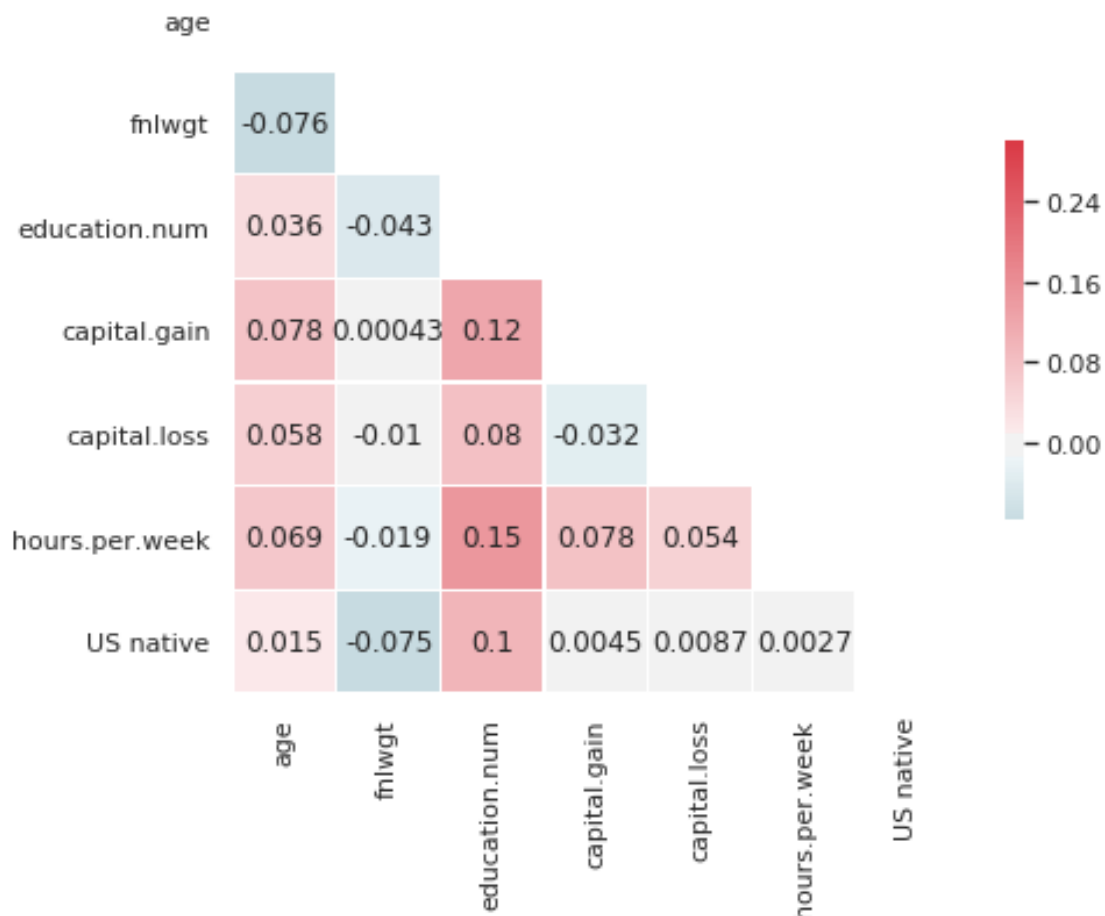
# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(7, 6))

# Generate a custom diverging colormap
cmap = sns.diverging_palette(220, 10, as_cmap=True)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, vmax=.3, center=0,
            square=True, linewidths=.5, annot=True, cbar_kws={"shrink": .5})

Out[30]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe274a74710>

```



---

## 4 Preparing data

```
In [31]: df['income'] = pd.get_dummies(df['income'], prefix='income', drop_first=True)
```

```
In [32]: y = df.income
         df = df.drop(columns=['income'])
```

```
In [33]: print(f'Ratio above 50k: {y.sum()/len(y)*100:.2f}%')
```

Ratio above 50k: 24.09%

```
In [34]: #cat_columns = ['workclass', 'education', 'marital-status', 'occupation', 'relationships']
```

```
In [35]: #df_clean['sex'] = df_clean['sex'].str.replace('Female', '0').str.replace('Male', '1')
```

```
In [36]: df.head()
```

```
Out [36]:
```

	age	workclass	fnlwgt	education	education.num	marital.status	\
0	90	?	77053	HS-grad	9	Widowed	
1	82	Private	132870	HS-grad	9	Widowed	
2	66	?	186061	Some-college	10	Widowed	
3	54	Private	140359	7th-8th	4	Divorced	
4	41	Private	264663	Some-college	10	Separated	

	occupation	relationship	race	sex	capital.gain	\
0	?	Not-in-family	White	Female	0	
1	Exec-managerial	Not-in-family	White	Female	0	
2	?	Unmarried	Black	Female	0	
3	Machine-op-inspct	Unmarried	White	Female	0	
4	Prof-specialty	Own-child	White	Female	0	

	capital.loss	hours.per.week	native.country	US native
0	4356	40	United-States	1
1	4356	18	United-States	1
2	4356	40	United-States	1
3	3900	40	United-States	1
4	3900	40	United-States	1

```
In [37]: cols = list(df.columns)
         cols
```

```
Out [37]: ['age',
           'workclass',
           'fnlwgt',
           'education',
           'education.num',
           'marital.status',
           'occupation',
           'relationship',
           'race',
           'sex',
           'capital.gain',
           'capital.loss',
           'hours.per.week',
           'native.country',
           'US native']
```

```
In [38]: selected_feat = cols.copy()
         selected_feat.remove('US native')
         selected_feat
```

```
Out [38]: ['age',
           'workclass',
           'fnlwgt',
           'education',
           'education.num',
           'marital.status',
           'occupation',
           'relationship',
           'race',
           'sex',
           'capital.gain',
           'capital.loss',
           'hours.per.week',
           'native.country']
```

```
In [42]: df_final = df[selected_feat]
```

```
In [43]: cat_feat = df_final.select_dtypes(include=['object']).columns
         X = pd.get_dummies(df_final[cat_feat], drop_first=True)
```

```
In [44]: #X = pd.concat([df_final[continuous_columns], df_dummies], axis=1)
```

```
In [45]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

---

## 5 Model training and predictions

Choose carefully your features, you can even create new ones, and focus on training Decision Trees and Random Forests.

Improve your models by finding the optimal hyperparameters that minimize your models' error.

```
In [46]: def print_score(model, name):  
         model.fit(X_train, y_train)  
         print('Accuracy score of the', name, f': on train = {model.score(X_train, y_train)}
```

## 5.1 Baseline LogisticRegression

```
In [47]: print_score(LogisticRegression(), 'LogisticReg')
```

Accuracy score of the LogisticReg : on train = 83.61%, on test = 81.78%

## 5.2 Decision Tree

```
In [48]: print_score(DecisionTreeClassifier(), 'DecisionTreeClf')
```

Accuracy score of the DecisionTreeClf : on train = 87.14%, on test = 80.19%

## 5.3 Random Forest

```
In [49]: rf = RandomForestClassifier().fit(X_train, y_train)  
         print(f'Accuracy score of the RandomForrest: on train = {rf.score(X_train, y_train)*100}
```

Accuracy score of the RandomForrest: on train = 86.83%, on test = 80.64%

## 5.4 ExtraTreesClassifier

```
In [50]: # fit an Extra Tree model to the data  
         print_score(DecisionTreeClassifier(), 'ExtraTreesClf')
```

Accuracy score of the ExtraTreesClf : on train = 87.14%, on test = 80.26%

## 5.5 Tuned model

```
In [53]: rfc = RandomForestClassifier()  
         param_grid = {  
             'n_estimators': [50, 100, 150, 200, 250],  
             'max_features': [1, 2, 3, 4, 5],  
             'max_depth' : [4, 6, 8]  
         }  
  
In [54]: rfc_cv = GridSearchCV(estimator=rfc, param_grid=param_grid, cv=5)  
         rfc_cv.fit(X_train, y_train)
```

```
Out [54]: GridSearchCV(cv=5, error_score='raise-deprecating',
                      estimator=RandomForestClassifier(bootstrap=True, class_weight=None, criterion=
                        max_depth=None, max_features='auto', max_leaf_nodes=None,
                        min_impurity_decrease=0.0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0.0, n_estimators='warn', n_jobs=None,
                        oob_score=False, random_state=None, verbose=0,
                        warm_start=False),
                      fit_params=None, iid='warn', n_jobs=None,
                      param_grid={'n_estimators': [50, 100, 150, 200, 250], 'max_features': [1, 2, 3
                        pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
                        scoring=None, verbose=0)
```

```
In [55]: rfc_cv.best_params_
```

```
Out [55]: {'max_depth': 8, 'max_features': 5, 'n_estimators': 250}
```

```
In [57]: rfc_best = RandomForestClassifier(max_depth=8, max_features=5, n_estimators=250).fit(X_train, y_train)
print(f'Accuracy score of the RandomForrest: on train = {rfc_best.score(X_train, y_train)}')
```

```
Accuracy score of the RandomForrest: on train = 80.54%, on test = 79.78%
```

## 6 Profiling

Let's find clear insights on the profiles of the people that make more than USD 50K a year. Which features seem to be the most correlated with this phenomenon.

### 6.1 Based on the rf model

```
In [58]: # indexes of columns which are the most important
np.argsort(rf.feature_importances_)[-16:]
```

```
Out [58]: array([91,  5, 22, 36, 21,  3, 18, 45, 43, 19, 38, 16, 52, 32, 26, 24])
```

```
In [59]: # most important features
[list(X.columns)[i] for i in np.argsort(rf.feature_importances_)[-16:]][::-1]
```

```
Out [59]: ['marital.status_Married-civ-spouse',
           'marital.status_Never-married',
           'occupation_Exec-managerial',
           'sex_Male',
           'education_Bachelors',
           'occupation_Prof-specialty',
           'education_Masters',
           'relationship_Not-in-family',
           'relationship_Own-child',
           'education_HS-grad',
           'workclass_Private',
```



```

'education_Prof-school',
'occupation_Other-service',
'education_Some-college',
'workclass_Self-emp-not-inc',
'native.country_United-States']

```

In [60]: *# Feature importances*

```

features = X.columns
importances = rf.feature_importances_
indices = np.argsort(importances)[::-1]
num_features = len(importances)

```

*# Plot the feature importances of the tree*

```

plt.figure(figsize=(16, 4))
plt.title("Feature importances")
plt.bar(range(num_features), importances[indices], color="g", align="center")
plt.xticks(range(num_features), [features[i] for i in indices], rotation='45')
plt.xlim([-1, num_features])
plt.show()

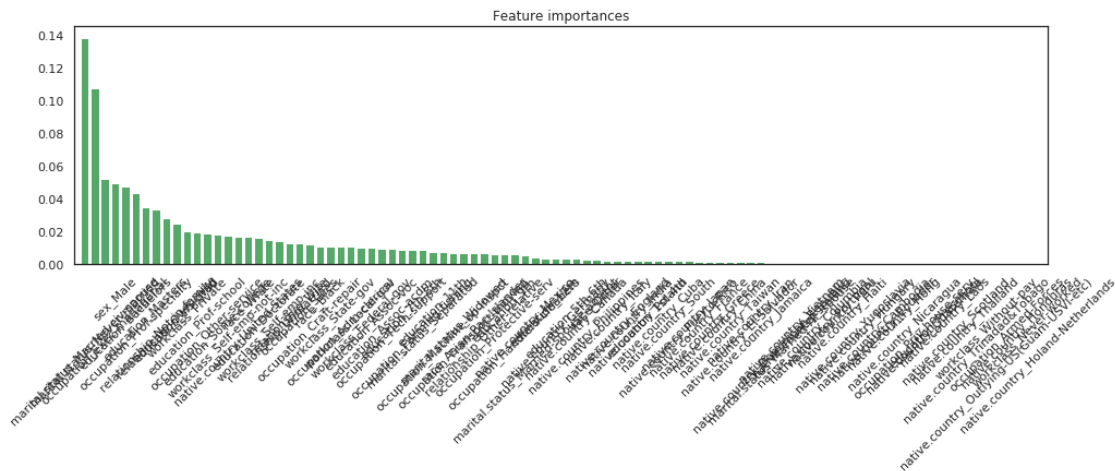
```

*# Print values*

```

for i in indices:
    print ("{0} - {1:.3f}".format(features[i], importances[i]))

```



```

marital.status_Married-civ-spouse - 0.139
marital.status_Never-married - 0.108
occupation_Exec-managerial - 0.052
sex_Male - 0.050
education_Bachelors - 0.047
occupation_Prof-specialty - 0.043
education_Masters - 0.034

```

relationship\_Not-in-family - 0.033  
relationship\_Own-child - 0.028  
education\_HS-grad - 0.025  
workclass\_Private - 0.020  
education\_Prof-school - 0.019  
occupation\_Other-service - 0.019  
education\_Some-college - 0.018  
workclass\_Self-emp-not-inc - 0.017  
native.country\_United-States - 0.017  
race\_White - 0.017  
education\_Doctorate - 0.016  
workclass\_Self-emp-inc - 0.015  
relationship\_Unmarried - 0.014  
relationship\_Wife - 0.013  
occupation\_Sales - 0.013  
occupation\_Craft-repair - 0.012  
race\_Black - 0.011  
workclass\_State-gov - 0.011  
occupation\_Adm-clerical - 0.011  
workclass\_Local-gov - 0.011  
workclass\_Federal-gov - 0.010  
education\_Assoc-voc - 0.010  
education\_Assoc-acdm - 0.010  
occupation\_Tech-support - 0.010  
education\_7th-8th - 0.009  
occupation\_Farming-fishing - 0.009  
marital.status\_Separated - 0.008  
education\_11th - 0.007  
occupation\_Machine-op-inspct - 0.007  
marital.status\_Widowed - 0.007  
occupation\_Transport-moving - 0.007  
race\_Asian-Pac-Islander - 0.007  
relationship\_Other-relative - 0.007  
occupation\_Protective-serv - 0.006  
education\_9th - 0.006  
occupation\_Handlers-cleaners - 0.006  
native.country\_Mexico - 0.005  
marital.status\_Married-spouse-absent - 0.004  
education\_12th - 0.004  
native.country\_Germany - 0.003  
education\_5th-6th - 0.003  
native.country\_Canada - 0.003  
native.country\_Philippines - 0.003  
race\_Other - 0.003  
native.country\_Italy - 0.002  
native.country\_England - 0.002  
native.country\_India - 0.002  
native.country\_Poland - 0.002

```

education_1st-4th - 0.002
native.country_Cuba - 0.002
native.country_South - 0.002
native.country_Puerto-Rico - 0.002
native.country_Japan - 0.002
native.country_France - 0.001
native.country_Greece - 0.001
native.country_China - 0.001
native.country_Taiwan - 0.001
native.country_El-Salvador - 0.001
native.country_Iran - 0.001
native.country_Jamaica - 0.001
native.country_Dominican-Republic - 0.001
marital.status_Married-AF-spouse - 0.001
native.country_Vietnam - 0.001
native.country_Ireland - 0.001
native.country_Columbia - 0.001
native.country_Portugal - 0.001
native.country_Peru - 0.001
native.country_Haiti - 0.001
native.country_Yugoslavia - 0.001
native.country_Hungary - 0.001
native.country_Cambodia - 0.000
native.country_Ecuador - 0.000
native.country_Hong - 0.000
native.country_Nicaragua - 0.000
occupation_Priv-house-serv - 0.000
native.country_Guatemala - 0.000
education_Preschool - 0.000
native.country_Laos - 0.000
native.country_Scotland - 0.000
native.country_Thailand - 0.000
native.country_Trinidad&Tobago - 0.000
workclass_Without-pay - 0.000
native.country_Outlying-US(Guam-USVI-etc) - 0.000
occupation_Armed-Forces - 0.000
native.country_Honduras - 0.000
workclass_Never-worked - 0.000
native.country_Holand-Netherlands - 0.000

```

```

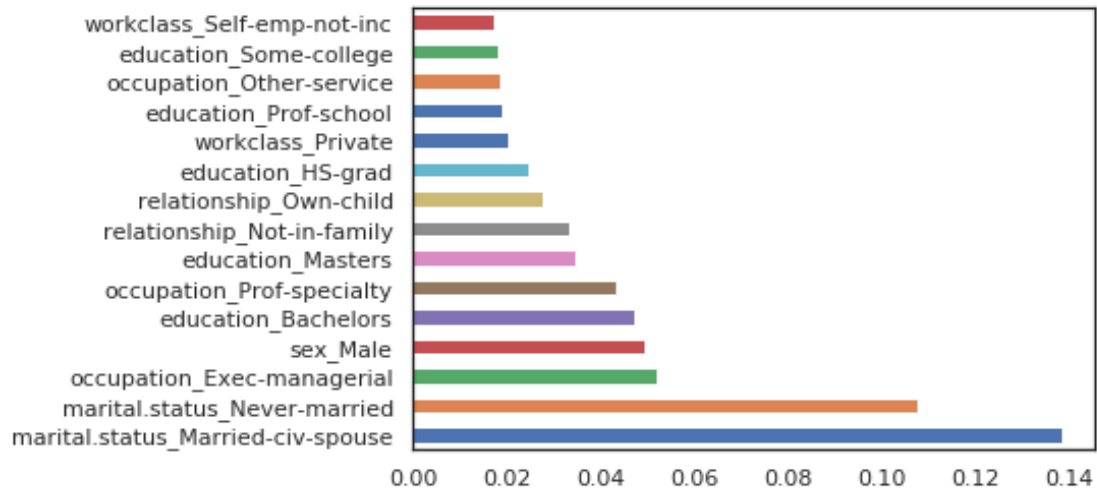
In [61]: (pd.Series(rf.feature_importances_, index=X_train.columns)
          .nlargest(15)
          .plot(kind='barh'))

```

```

Out[61]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe2780ecef0>

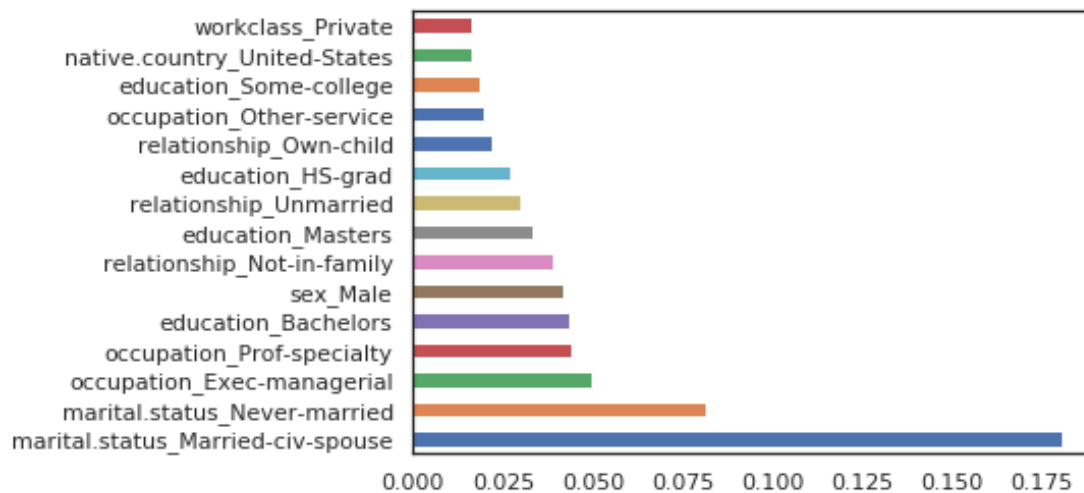
```



## 6.2 Based on the ExtraTree model

```
In [62]: extree = ExtraTreesClassifier().fit(X_train, y_train)
          (pd.Series(extree.feature_importances_, index=X_train.columns)
           .nlargest(15)
           .plot(kind='barh'))
```

```
Out [62]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe277f0f7f0>
```



The same features come first.