

Home_credit_default_risk

June 24, 2019

1 Home Credit Default Risk

Predict how capable each applicant is of repaying a loan.

Photo [Breno Assis](#)

1.1 Context

This [challenge](#) was proposed by **Home Credit Group**.

Many people struggle to get loans due to insufficient or non-existent credit histories. And, unfortunately, this population is often taken advantage of by untrustworthy lenders.

Home Credit strives to broaden financial inclusion for the unbanked population by providing a positive and safe borrowing experience. In order to make sure this underserved population has a positive loan experience, Home Credit makes use of a variety of alternative data--including telco and transactional information--to predict their clients' repayment abilities.

While Home Credit is currently using various statistical and machine learning methods to make these predictions, they're challenging Kagglers to help them unlock the full potential of their data. Doing so will ensure that clients capable of repayment are not rejected and that loans are given with a principal, maturity, and repayment calendar that will empower their clients to be successful.

1.2 Goal

Use historical loan application data to **predict whether or not an applicant will be able to repay a loan**. This is a standard **supervised classification task**.

Submissions are evaluated on **area under the ROC curve** between the predicted probability and the observed target.

1.3 Type of ML

- Supervised: Labels are included in the training data
- Binary classification: target has only two values 0 (will repay loan on time), 1 (will have difficulty repaying loan)

1.4 Guidelines

- Download and load the data
- Sample the data in order to work on a smaller subset at first
- Explore the data, creating functions for cleaning it

- Split your data into features & labels ; training & testing
 - Train different models and compare performance
 - Train on your entire dataset, run predictions on your entire dataset and submit your results!
 - Iterate
-

2 Data

2.1 First insight

In [1]: *# usual data science stack in python*

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
import os
```

```
print(os.listdir("../input/"))
```

Any results you write to the current directory are saved as output.

```
['application_train.csv', 'home_credit.png', 'ROC-curve.png', 'HomeCredit_columns_description_1.png']
```

In [2]: *# imports of need modules in sklearn*

```
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import cross_val_score
```

```
from sklearn.ensemble import RandomForestClassifier
```

In [3]: *import lightgbm as lgb*

```
import xgboost as xgb
```

In [4]: *# set options in this notebook*

```
pd.set_option('display.max_columns', 300)
```

```
import warnings
```

```
warnings.simplefilter(action='ignore', category=FutureWarning)
```

In [5]: *path_train = os.path.join('../', 'input', 'application_train.csv')*

```
path_test = os.path.join('../', 'input', 'application_test.csv')
```

In [6]: *# load main datasets*

```
app_train, app_test = pd.read_csv(path_train), pd.read_csv(path_test)
```

In [7]: *# 1st insight*

```
app_train.tail()
```

Out [7]:

	SK_ID_CURR	TARGET	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	\
307506	456251	0	Cash loans	M	N	
307507	456252	0	Cash loans	F	N	
307508	456253	0	Cash loans	F	N	
307509	456254	1	Cash loans	F	N	
307510	456255	0	Cash loans	F	N	

	FLAG_OWN_REALTY	CNT_CHILDREN	AMT_INCOME_TOTAL	AMT_CREDIT	\
307506	N	0	157500.0	254700.0	
307507	Y	0	72000.0	269550.0	
307508	Y	0	153000.0	677664.0	
307509	Y	0	171000.0	370107.0	
307510	N	0	157500.0	675000.0	

	AMT_ANNUITY	AMT_GOODS_PRICE	NAME_TYPE_SUITE	NAME_INCOME_TYPE	\
307506	27558.0	225000.0	Unaccompanied	Working	
307507	12001.5	225000.0	Unaccompanied	Pensioner	
307508	29979.0	585000.0	Unaccompanied	Working	
307509	20205.0	319500.0	Unaccompanied	Commercial associate	
307510	49117.5	675000.0	Unaccompanied	Commercial associate	

	NAME_EDUCATION_TYPE	NAME_FAMILY_STATUS	NAME_HOUSING_TYPE	\
307506	Secondary / secondary special	Separated	With parents	
307507	Secondary / secondary special	Widow	House / apartment	
307508	Higher education	Separated	House / apartment	
307509	Secondary / secondary special	Married	House / apartment	
307510	Higher education	Married	House / apartment	

	REGION_POPULATION_RELATIVE	DAYS_BIRTH	DAYS_EMPLOYED	\
307506	0.032561	-9327	-236	
307507	0.025164	-20775	365243	
307508	0.005002	-14966	-7921	
307509	0.005313	-11961	-4786	
307510	0.046220	-16856	-1262	

	DAYS_REGISTRATION	DAYS_ID_PUBLISH	OWN_CAR_AGE	FLAG_MOBIL	\
307506	-8456.0	-1982	NaN	1	
307507	-4388.0	-4090	NaN	1	
307508	-6737.0	-5150	NaN	1	
307509	-2562.0	-931	NaN	1	
307510	-5128.0	-410	NaN	1	

	FLAG_EMP_PHONE	FLAG_WORK_PHONE	FLAG_CONT_MOBILE	FLAG_PHONE	\
307506	1	0	1	0	
307507	0	0	1	1	
307508	1	0	1	0	
307509	1	0	1	0	
307510	1	1	1	1	

	FLAG_EMAIL	OCCUPATION_TYPE	CNT_FAM_MEMBERS	REGION_RATING_CLIENT	\
307506	0	Sales staff	1.0		1
307507	0	NaN	1.0		2
307508	1	Managers	1.0		3
307509	0	Laborers	2.0		2
307510	0	Laborers	2.0		1

	REGION_RATING_CLIENT_W_CITY	WEEKDAY_APPR_PROCESS_START	\
307506		1 THURSDAY	
307507		2 MONDAY	
307508		3 THURSDAY	
307509		2 WEDNESDAY	
307510		1 THURSDAY	

	HOURLY_APPR_PROCESS_START	REG_REGION_NOT_LIVE_REGION	\
307506	15	0	
307507	8	0	
307508	9	0	
307509	9	0	
307510	20	0	

	REG_REGION_NOT_WORK_REGION	LIVE_REGION_NOT_WORK_REGION	\
307506	0	0	
307507	0	0	
307508	0	0	
307509	0	0	
307510	0	0	

	REG_CITY_NOT_LIVE_CITY	REG_CITY_NOT_WORK_CITY	\
307506	0	0	
307507	0	0	
307508	0	1	
307509	1	1	
307510	0	1	

	LIVE_CITY_NOT_WORK_CITY	ORGANIZATION_TYPE	EXT_SOURCE_1	\
307506	0	Services	0.145570	
307507	0	XNA	NaN	
307508	1	School	0.744026	
307509	0	Business Entity Type 1	NaN	
307510	1	Business Entity Type 3	0.734460	

	EXT_SOURCE_2	EXT_SOURCE_3	APARTMENTS_AVG	BASEMENTAREA_AVG	\
307506	0.681632	NaN	0.2021	0.0887	
307507	0.115992	NaN	0.0247	0.0435	
307508	0.535722	0.218859	0.1031	0.0862	
307509	0.514163	0.661024	0.0124	NaN	

307510	0.708569	0.113922	0.0742	0.0526
--------	----------	----------	--------	--------

	YEARS_BEGINEXPLUATATION_AVG	YEARS_BUILD_AVG	COMMONAREA_AVG	\
307506	0.9876	0.8300	0.0202	
307507	0.9727	0.6260	0.0022	
307508	0.9816	0.7484	0.0123	
307509	0.9771	NaN	NaN	
307510	0.9881	NaN	0.0176	

	ELEVATORS_AVG	ENTRANCES_AVG	FLOORSMAX_AVG	FLOORSMIN_AVG	\
307506	0.22	0.1034	0.6042	0.2708	
307507	0.00	0.1034	0.0833	0.1250	
307508	0.00	0.2069	0.1667	0.2083	
307509	NaN	0.0690	0.0417	NaN	
307510	0.08	0.0690	0.3750	NaN	

	LANDAREA_AVG	LIVINGAPARTMENTS_AVG	LIVINGAREA_AVG	\
307506	0.0594	0.1484	0.1965	
307507	0.0579	0.0202	0.0257	
307508	NaN	0.0841	0.9279	
307509	NaN	NaN	0.0061	
307510	NaN	NaN	0.0791	

	NONLIVINGAPARTMENTS_AVG	NONLIVINGAREA_AVG	APARTMENTS_MODE	\
307506	0.0753	0.1095	0.1008	
307507	0.0000	0.0000	0.0252	
307508	0.0000	0.0000	0.1050	
307509	NaN	NaN	0.0126	
307510	NaN	0.0000	0.0756	

	BASEMENTAREA_MODE	YEARS_BEGINEXPLUATATION_MODE	YEARS_BUILD_MODE	\
307506	0.0172	0.9782	0.7125	
307507	0.0451	0.9727	0.6406	
307508	0.0894	0.9816	0.7583	
307509	NaN	0.9772	NaN	
307510	0.0546	0.9881	NaN	

	COMMONAREA_MODE	ELEVATORS_MODE	ENTRANCES_MODE	FLOORSMAX_MODE	\
307506	0.0172	0.0806	0.0345	0.4583	
307507	0.0022	0.0000	0.1034	0.0833	
307508	0.0124	0.0000	0.2069	0.1667	
307509	NaN	NaN	0.0690	0.0417	
307510	0.0178	0.0806	0.0690	0.3750	

	FLOORSMIN_MODE	LANDAREA_MODE	LIVINGAPARTMENTS_MODE	LIVINGAREA_MODE	\
307506	0.0417	0.0094	0.0882	0.0853	
307507	0.1250	0.0592	0.0220	0.0267	
307508	0.2083	NaN	0.0918	0.9667	

307509	NaN	NaN	NaN	0.0063
307510	NaN	NaN	NaN	0.0824

	NONLIVINGAPARTMENTS_MODE	NONLIVINGAREA_MODE	APARTMENTS_MEDI	\
307506	0.0	0.0125	0.2040	
307507	0.0	0.0000	0.0250	
307508	0.0	0.0000	0.1041	
307509	NaN	NaN	0.0125	
307510	NaN	0.0000	0.0749	

	BASEMENTAREA_MEDI	YEARS_BEGINEXPLUATATION_MEDI	YEARS_BUILD_MEDI	\
307506	0.0887	0.9876	0.8323	
307507	0.0435	0.9727	0.6310	
307508	0.0862	0.9816	0.7518	
307509	NaN	0.9771	NaN	
307510	0.0526	0.9881	NaN	

	COMMONAREA_MEDI	ELEVATORS_MEDI	ENTRANCES_MEDI	FLOORSMAX_MEDI	\
307506	0.0203	0.22	0.1034	0.6042	
307507	0.0022	0.00	0.1034	0.0833	
307508	0.0124	0.00	0.2069	0.1667	
307509	NaN	NaN	0.0690	0.0417	
307510	0.0177	0.08	0.0690	0.3750	

	FLOORSMIN_MEDI	LANDAREA_MEDI	LIVINGAPARTMENTS_MEDI	LIVINGAREA_MEDI	\
307506	0.2708	0.0605	0.1509	0.2001	
307507	0.1250	0.0589	0.0205	0.0261	
307508	0.2083	NaN	0.0855	0.9445	
307509	NaN	NaN	NaN	0.0062	
307510	NaN	NaN	NaN	0.0805	

	NONLIVINGAPARTMENTS_MEDI	NONLIVINGAREA_MEDI	FONDKAPREMONT_MODE	\
307506	0.0757	0.1118	reg oper account	
307507	0.0000	0.0000	reg oper account	
307508	0.0000	0.0000	reg oper account	
307509	NaN	NaN	NaN	
307510	NaN	0.0000	NaN	

	HOUSETYPE_MODE	TOTALAREA_MODE	WALLSMATERIAL_MODE	EMERGENCYSTATE_MODE	\
307506	block of flats	0.2898	Stone, brick	No	
307507	block of flats	0.0214	Stone, brick	No	
307508	block of flats	0.7970	Panel	No	
307509	block of flats	0.0086	Stone, brick	No	
307510	block of flats	0.0718	Panel	No	

	OBS_30_CNT_SOCIAL_CIRCLE	DEF_30_CNT_SOCIAL_CIRCLE	\
307506	0.0	0.0	
307507	0.0	0.0	

307508	6.0	0.0
307509	0.0	0.0
307510	0.0	0.0

	OBS_60_CNT_SOCIAL_CIRCLE	DEF_60_CNT_SOCIAL_CIRCLE	\
307506	0.0	0.0	
307507	0.0	0.0	
307508	6.0	0.0	
307509	0.0	0.0	
307510	0.0	0.0	

	DAYS_LAST_PHONE_CHANGE	FLAG_DOCUMENT_2	FLAG_DOCUMENT_3	\
307506	-273.0	0	0	
307507	0.0	0	1	
307508	-1909.0	0	1	
307509	-322.0	0	1	
307510	-787.0	0	1	

	FLAG_DOCUMENT_4	FLAG_DOCUMENT_5	FLAG_DOCUMENT_6	FLAG_DOCUMENT_7	\
307506	0	0	0	0	
307507	0	0	0	0	
307508	0	0	0	0	
307509	0	0	0	0	
307510	0	0	0	0	

	FLAG_DOCUMENT_8	FLAG_DOCUMENT_9	FLAG_DOCUMENT_10	FLAG_DOCUMENT_11	\
307506	1	0	0	0	
307507	0	0	0	0	
307508	0	0	0	0	
307509	0	0	0	0	
307510	0	0	0	0	

	FLAG_DOCUMENT_12	FLAG_DOCUMENT_13	FLAG_DOCUMENT_14	\
307506	0	0	0	
307507	0	0	0	
307508	0	0	0	
307509	0	0	0	
307510	0	0	0	

	FLAG_DOCUMENT_15	FLAG_DOCUMENT_16	FLAG_DOCUMENT_17	\
307506	0	0	0	
307507	0	0	0	
307508	0	0	0	
307509	0	0	0	
307510	0	0	0	

	FLAG_DOCUMENT_18	FLAG_DOCUMENT_19	FLAG_DOCUMENT_20	\
307506	0	0	0	

307507	0	0	0
307508	0	0	0
307509	0	0	0
307510	0	0	0

	FLAG_DOCUMENT_21	AMT_REQ_CREDIT_BUREAU_HOUR	\
307506	0	NaN	
307507	0	NaN	
307508	0	1.0	
307509	0	0.0	
307510	0	0.0	

	AMT_REQ_CREDIT_BUREAU_DAY	AMT_REQ_CREDIT_BUREAU_WEEK	\
307506	NaN	NaN	
307507	NaN	NaN	
307508	0.0	0.0	
307509	0.0	0.0	
307510	0.0	0.0	

	AMT_REQ_CREDIT_BUREAU_MON	AMT_REQ_CREDIT_BUREAU_QRT	\
307506	NaN	NaN	
307507	NaN	NaN	
307508	1.0	0.0	
307509	0.0	0.0	
307510	2.0	0.0	

	AMT_REQ_CREDIT_BUREAU_YEAR
307506	NaN
307507	NaN
307508	1.0
307509	0.0
307510	1.0

In [8]: `app_train.shape, app_test.shape`

Out [8]: ((307511, 122), (48744, 121))

Of course, the column named 'TARGET' is not in the test dataset.

2.2 Content of each table and links between tables

- **application_{train | test}.csv**
 - This is the main table, broken into two files for Train (with TARGET) and Test (without TARGET).
 - Static data for all applications. One row represents one loan in our data sample.
- **bureau.csv**
 - All client's previous credits provided by other financial institutions that were reported to Credit Bureau (for clients who have a loan in our sample).

- For every loan in our sample, there are as many rows as number of credits the client had in Credit Bureau before the application date.
 - **bureau_balance.csv**
 - Monthly balances of previous credits in Credit Bureau.
 - This table has one row for each month of history of every previous credit reported to Credit Bureau – i.e the table has (#loans in sample * # of relative previous credits * # of months where we have some history observable for the previous credits) rows.
 - **POS_CASH_balance.csv**
 - Monthly balance snapshots of previous POS (point of sales) and cash loans that the applicant had with Home Credit.
 - This table has one row for each month of history of every previous credit in Home Credit (consumer credit and cash loans) related to loans in our sample – i.e. the table has (#loans in sample * # of relative previous credits * # of months in which we have some history observable for the previous credits) rows.
 - **credit_card_balance.csv**
 - Monthly balance snapshots of previous credit cards that the applicant has with Home Credit.
 - This table has one row for each month of history of every previous credit in Home Credit (consumer credit and cash loans) related to loans in our sample – i.e. the table has (#loans in sample * # of relative previous credit cards * # of months where we have some history observable for the previous credit card) rows.
 - **previous_application.csv**
 - All previous applications for Home Credit loans of clients who have loans in our sample.
 - There is one row for each previous application related to loans in our data sample.
 - **installments_payments.csv**
 - Repayment history for the previously disbursed credits in Home Credit related to the loans in our sample.
 - There is a) one row for every payment that was made plus b) one row each for missed payment.
 - One row is equivalent to one payment of one installment OR one installment corresponding to one payment of one previous Home Credit credit related to loans in our sample.
 - **HomeCredit_columns_description.csv**
 - This file contains descriptions for the columns in the various data files.
-

3 Exploratory Data Analysis

3.1 Distribution of the Target Column


In [9]: `app_train.TARGET.value_counts()`

```
Out[9]: 0    282686
        1     24825
        Name: TARGET, dtype: int64
```

```
In [10]: print(f'percentage of clients with payment difficulties: {app_train.TARGET.sum() / app_train.TARGET.count()}')
```

percentage of clients with payment difficulties: 8.07%

```
In [11]: plt.title('Distribution of the Target Column - 1 - client with payment difficulties / 0 - all other cases')
        sns.countplot(x=app_train.TARGET, data=app_train)
        plt.show()
```



This is an [imbalanced class problem](#). There are far more repaid loans than loans that were not repaid. It is important to weight the classes by their representation in the data to reflect this imbalance.

3.2 Column Types

```
In [12]: app_train.dtypes.value_counts()
```

```
Out[12]: float64    65
         int64     41
         object     16
         dtype: int64
```

int64 and float64 are numeric variables which can correspond to discrete or continuous features. Whereas object columns contain strings and are categorical features.

```
In [13]: # Number of unique classes in each object column
        app_train.select_dtypes('object').apply(pd.Series.nunique, axis=0)
```

```

Out[13]: NAME_CONTRACT_TYPE          2
         CODE_GENDER                  3
         FLAG_OWN_CAR                 2
         FLAG_OWN_REALTY              2
         NAME_TYPE_SUITE              7
         NAME_INCOME_TYPE             8
         NAME_EDUCATION_TYPE          5
         NAME_FAMILY_STATUS           6
         NAME_HOUSING_TYPE            6
         OCCUPATION_TYPE             18
         WEEKDAY_APPR_PROCESS_START   7
         ORGANIZATION_TYPE           58
         FONDKAPREMONT_MODE           4
         HOUSETYPE_MODE               3
         WALLSMATERIAL_MODE           7
         EMERGENCYSTATE_MODE          2
         dtype: int64

```

3.3 Missing Values

```

In [14]: # Function to calculate missing values by column# Funct // credits Will Koehrsen
def missing_values_table(df):
    # Total missing values
    mis_val = df.isnull().sum()

    # Percentage of missing values
    mis_val_percent = 100 * df.isnull().sum() / len(df)

    # Make a table with the results
    mis_val_table = pd.concat([mis_val, mis_val_percent], axis=1)

    # Rename the columns
    mis_val_table_ren_columns = mis_val_table.rename(
        columns = {0 : 'Missing Values', 1 : '% of Total Values'})

    # Sort the table by percentage of missing descending
    mis_val_table_ren_columns = mis_val_table_ren_columns[
        mis_val_table_ren_columns.iloc[:,1] != 0].sort_values(
        '% of Total Values', ascending=False).round(1)

    # Print some summary information
    print ("Your selected dataframe has " + str(df.shape[1]) + " columns.\n"
          "There are " + str(mis_val_table_ren_columns.shape[0]) +
          " columns that have missing values.")

    # Return the dataframe with missing information
    return mis_val_table_ren_columns

```

```
In [15]: # Missing values statistics
missing_values = missing_values_table(app_train)
missing_values.head(10)
```

Your selected dataframe has 122 columns.
There are 67 columns that have missing values.

```
Out [15]:
```

	Missing Values	% of Total Values
COMMONAREA_MEDI	214865	69.9
COMMONAREA_AVG	214865	69.9
COMMONAREA_MODE	214865	69.9
NONLIVINGAPARTMENTS_MEDI	213514	69.4
NONLIVINGAPARTMENTS_MODE	213514	69.4
NONLIVINGAPARTMENTS_AVG	213514	69.4
FONDKAPREMONT_MODE	210295	68.4
LIVINGAPARTMENTS_MODE	210199	68.4
LIVINGAPARTMENTS_MEDI	210199	68.4
LIVINGAPARTMENTS_AVG	210199	68.4

From here we have 2 options : * Use models such as XGBoost that can handle missing values
* Or drop columns with a high percentage of missing values, and fill in other columns with a low percentage It is not possible to know ahead of time if these columns will be helpful or not. My choice here is to drop them. Later if we need a more accurate score, we'll change the way to proceed.

3.3.1 Dropping columns with a high ratio of missing values

```
In [16]: # cols_to_drop = list((app_train.isnull().sum() > 75000).index)
cols_to_drop = [c for c in app_train.columns if app_train[c].isnull().sum() > 75000]

In [17]: app_train, app_test = app_train.drop(cols_to_drop, axis=1), app_test.drop(cols_to_drop, axis=1)
app_test.isnull().sum().sort_values(ascending=False).head(10)
```

```
Out [17]: EXT_SOURCE_3      8668
AMT_REQ_CREDIT_BUREAU_YEAR  6049
AMT_REQ_CREDIT_BUREAU_MON   6049
AMT_REQ_CREDIT_BUREAU_WEEK  6049
AMT_REQ_CREDIT_BUREAU_DAY   6049
AMT_REQ_CREDIT_BUREAU_HOUR  6049
AMT_REQ_CREDIT_BUREAU_QRT   6049
NAME_TYPE_SUITE             911
DEF_60_CNT_SOCIAL_CIRCLE     29
OBS_60_CNT_SOCIAL_CIRCLE     29
dtype: int64
```

3.3.2 Filling other missing values

```
In [18]: obj_cols = app_train.select_dtypes('object').columns
obj_cols
```

```
Out [18]: Index(['NAME_CONTRACT_TYPE', 'CODE_GENDER', 'FLAG_OWN_CAR', 'FLAG_OWN_REALTY',
               'NAME_TYPE_SUITE', 'NAME_INCOME_TYPE', 'NAME_EDUCATION_TYPE',
               'NAME_FAMILY_STATUS', 'NAME_HOUSING_TYPE', 'WEEKDAY_APPR_PROCESS_START',
               'ORGANIZATION_TYPE'],
              dtype='object')
```

```
In [19]: # filling string cols with 'Not specified'
app_train[obj_cols] = app_train[obj_cols].fillna('Not specified')
app_test[obj_cols] = app_test[obj_cols].fillna('Not specified')
```

```
In [20]: float_cols = app_train.select_dtypes('float').columns
float_cols
```

```
Out [20]: Index(['AMT_INCOME_TOTAL', 'AMT_CREDIT', 'AMT_ANNUITY', 'AMT_GOODS_PRICE',
               'REGION_POPULATION_RELATIVE', 'DAYS_REGISTRATION', 'CNT_FAM_MEMBERS',
               'EXT_SOURCE_2', 'EXT_SOURCE_3', 'OBS_30_CNT_SOCIAL_CIRCLE',
               'DEF_30_CNT_SOCIAL_CIRCLE', 'OBS_60_CNT_SOCIAL_CIRCLE',
               'DEF_60_CNT_SOCIAL_CIRCLE', 'DAYS_LAST_PHONE_CHANGE',
               'AMT_REQ_CREDIT_BUREAU_HOUR', 'AMT_REQ_CREDIT_BUREAU_DAY',
               'AMT_REQ_CREDIT_BUREAU_WEEK', 'AMT_REQ_CREDIT_BUREAU_MON',
               'AMT_REQ_CREDIT_BUREAU_QRT', 'AMT_REQ_CREDIT_BUREAU_YEAR'],
              dtype='object')
```

```
In [21]: # filling float values with median of train (not test)
app_train[float_cols] = app_train[float_cols].fillna(app_train[float_cols].median())
app_test[float_cols] = app_test[float_cols].fillna(app_test[float_cols].median())
```

```
In [22]: app_train.shape, app_test.shape
```

```
Out [22]: ((307511, 72), (48744, 71))
```

Let's check if there is still NaNs

```
In [23]: app_train.isnull().sum().sort_values(ascending=False).head()
```

```
Out [23]: AMT_REQ_CREDIT_BUREAU_YEAR    0
          AMT_REQ_CREDIT_BUREAU_QRT    0
          DAYS_REGISTRATION             0
          DAYS_ID_PUBLISH               0
          FLAG_MOBIL                    0
          dtype: int64
```

```
In [24]: app_test.isnull().sum().sort_values(ascending=False).head()
```

```
Out [24]: AMT_REQ_CREDIT_BUREAU_YEAR    0
          FLAG_EMAIL                     0
          DAYS_ID_PUBLISH               0
          FLAG_MOBIL                    0
          FLAG_EMP_PHONE                 0
          dtype: int64
```

```
In [25]: # Is there any duplicated rows ?
```

```
In [26]: app_train.duplicated().sum()
```

```
Out[26]: 0
```

```
In [27]: app_test.duplicated().sum()
```

```
Out[27]: 0
```

3.4 Categorical columns (type object)

```
In [28]: # Number of unique classes in each object column
```

```
app_train.select_dtypes('object').apply(pd.Series.nunique, axis = 0)
```

```
Out[28]: NAME_CONTRACT_TYPE      2
         CODE_GENDER             3
         FLAG_OWN_CAR            2
         FLAG_OWN_REALTY         2
         NAME_TYPE_SUITE          8
         NAME_INCOME_TYPE         8
         NAME_EDUCATION_TYPE       5
         NAME_FAMILY_STATUS       6
         NAME_HOUSING_TYPE         6
         WEEKDAY_APPR_PROCESS_START 7
         ORGANIZATION_TYPE        58
         dtype: int64
```


3.5 Dealing with anomalies

```
In [29]: app_train['DAYS_EMPLOYED'].describe()
```

```
Out[29]: count      307511.000000
         mean        63815.045904
         std         141275.766519
         min         -17912.000000
         25%         -2760.000000
         50%         -1213.000000
         75%         -289.000000
         max          365243.000000
         Name: DAYS_EMPLOYED, dtype: float64
```

The maximum value is abnormal (besides being positive). It corresponds to 1000 years...

```
In [30]: sns.distplot(app_train['DAYS_EMPLOYED'], kde=False);
         plt.show()
```



```
In [31]: print('The non-anomalies default on %0.2f%% of loans' % (100 * app_train[app_train['DAYS_EMPLOYED'] < 365243].DEFAULT_RATE.mean())
          print('The anomalies default on %0.2f%% of loans' % (100 * app_train[app_train['DAYS_EMPLOYED'] >= 365243].DEFAULT_RATE.mean())
          print('There are %d anomalous days of employment' % len(app_train[app_train['DAYS_EMPLOYED'] >= 365243]))
```

The non-anomalies default on 8.66% of loans

The anomalies default on 5.40% of loans

There are 55374 anomalous days of employment


It turns out that the anomalies have a lower rate of default.

The anomalous values seem to have some importance. Let's fill in the anomalous values with not a np.nan and then create a new boolean column indicating whether or not the value was anomalous.

```
In [32]: # Create an anomalous flag column
          app_train['DAYS_EMPLOYED_ANOM'] = app_train["DAYS_EMPLOYED"] == 365243

          # Replace the anomalous values with nan
          app_train['DAYS_EMPLOYED'].replace({365243: np.nan}, inplace = True)

          sns.distplot(app_train['DAYS_EMPLOYED'].dropna(), kde=False);
```



```
In [33]: app_test['DAYS_EMPLOYED_ANOM'] = app_test["DAYS_EMPLOYED"] == 365243
         app_test["DAYS_EMPLOYED"].replace({365243: np.nan}, inplace = True)

         print('There are %d anomalies in the test data out of %d entries' % (app_test["DAYS_EMPLOYED"].isna().sum(), app_test["DAYS_EMPLOYED"].count()))
```

There are 9274 anomalies in the test data out of 48744 entries

```
In [34]: # refilling float values with median of train (not test)

         app_train[float_cols] = app_train[float_cols].apply(pd.to_numeric, errors='coerce')
         app_train = app_train.fillna(app_train.median())

         app_test[float_cols] = app_test[float_cols].apply(pd.to_numeric, errors='coerce')
         app_test = app_test.fillna(app_test.median())
```

3.6 Correlations

The correlation coefficient is not the best method to represent "relevance" of a feature, but it gives us an idea of possible relationships within the data. Some general interpretations of the absolute value of the correlation coefficient are:

- 00-.19 "very weak"
- 20-.39 "weak"
- 40-.59 "moderate"

- 60-.79 “strong”
- 80-1.0 “very strong”

```
In [35]: correlations = app_train.corr()['TARGET'].sort_values()

print('Most Positive Correlations:\n', correlations.tail(10))
print('\n\nMost Negative Correlations:\n', correlations.head(10))
```

```
Most Positive Correlations:
REG_CITY_NOT_LIVE_CITY      0.044395
FLAG_EMP_PHONE              0.045982
REG_CITY_NOT_WORK_CITY      0.050994
DAYS_ID_PUBLISH             0.051457
DAYS_LAST_PHONE_CHANGE      0.055218
REGION_RATING_CLIENT        0.058899
REGION_RATING_CLIENT_W_CITY 0.060893
DAYS_EMPLOYED               0.063368
DAYS_BIRTH                  0.078239
TARGET                      1.000000
Name: TARGET, dtype: float64
```

```
Most Negative Correlations:
EXT_SOURCE_2                -0.160295
EXT_SOURCE_3                -0.155892
DAYS_EMPLOYED_ANOM          -0.045987
AMT_GOODS_PRICE             -0.039623
REGION_POPULATION_RELATIVE  -0.037227
AMT_CREDIT                  -0.030369
FLAG_DOCUMENT_6             -0.028602
HOUR_APPR_PROCESS_START     -0.024166
FLAG_PHONE                  -0.023806
AMT_REQ_CREDIT_BUREAU_MON   -0.014794
Name: TARGET, dtype: float64
```

```
In [36]: # Compute the correlation matrix
corr = app_train.corr()

# Generate a mask for the upper triangle
mask = np.zeros_like(corr, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True

# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(21, 19))

# Generate a custom diverging colormap
cmap = sns.diverging_palette(220, 10, as_cmap=True)
```

```
# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, vmax=.3, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5})
```

Out[36]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7b866f50f0>



3.6.1 Effect of Age on Repayment

```
In [37]: # Find the correlation of the positive days since birth and target
app_train['DAYS_BIRTH'] = abs(app_train['DAYS_BIRTH'])
app_train['DAYS_BIRTH'].corr(app_train['TARGET'])
```

Out[37]: -0.07823930830982712

There isn't any correlation between age and repayment

```
In [38]: plt.figure(figsize = (12, 6))
```

```
# KDE plot of loans that were repaid on time
```


```
sns.kdeplot(app_train.loc[app_train['TARGET'] == 0, 'DAYS_BIRTH'] / 365, label = 'target == 0')
```

```
# KDE plot of loans which were not repaid on time
```

```
sns.kdeplot(app_train.loc[app_train['TARGET'] == 1, 'DAYS_BIRTH'] / 365, label = 'target == 1')
```

```
# Labeling of plot
```

```
plt.xlabel('Age (years)'); plt.ylabel('Density'); plt.title('Distribution of Ages');
```



```
In [39]: # Age information into a separate dataframe
```

```
age_data = app_train[['TARGET', 'DAYS_BIRTH']]
```

```
age_data['YEARS_BIRTH'] = age_data['DAYS_BIRTH'] / 365
```

```
# Bin the age data
```

```
age_data['YEARS_BINNED'] = pd.cut(age_data['YEARS_BIRTH'], bins = np.linspace(20, 70, 10))  
age_data.head(10)
```

```
/home/sunflowa/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:3: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

This is separate from the ipykernel package so we can avoid doing imports until

```
/home/sunflowa/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:6: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.
```

```
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

```
Out[39]:
```

	TARGET	DAYS_BIRTH	YEARS_BIRTH	YEARS_BINNED
0	1	9461	25.920548	(25.0, 30.0]
1	0	16765	45.931507	(45.0, 50.0]
2	0	19046	52.180822	(50.0, 55.0]
3	0	19005	52.068493	(50.0, 55.0]
4	0	19932	54.608219	(50.0, 55.0]
5	0	16941	46.413699	(45.0, 50.0]
6	0	13778	37.747945	(35.0, 40.0]
7	0	18850	51.643836	(50.0, 55.0]
8	0	20099	55.065753	(55.0, 60.0]
9	0	14469	39.641096	(35.0, 40.0]

```
In [40]: # Group by the bin and calculate averages
age_groups = age_data.groupby('YEARS_BINNED').mean()
age_groups
```

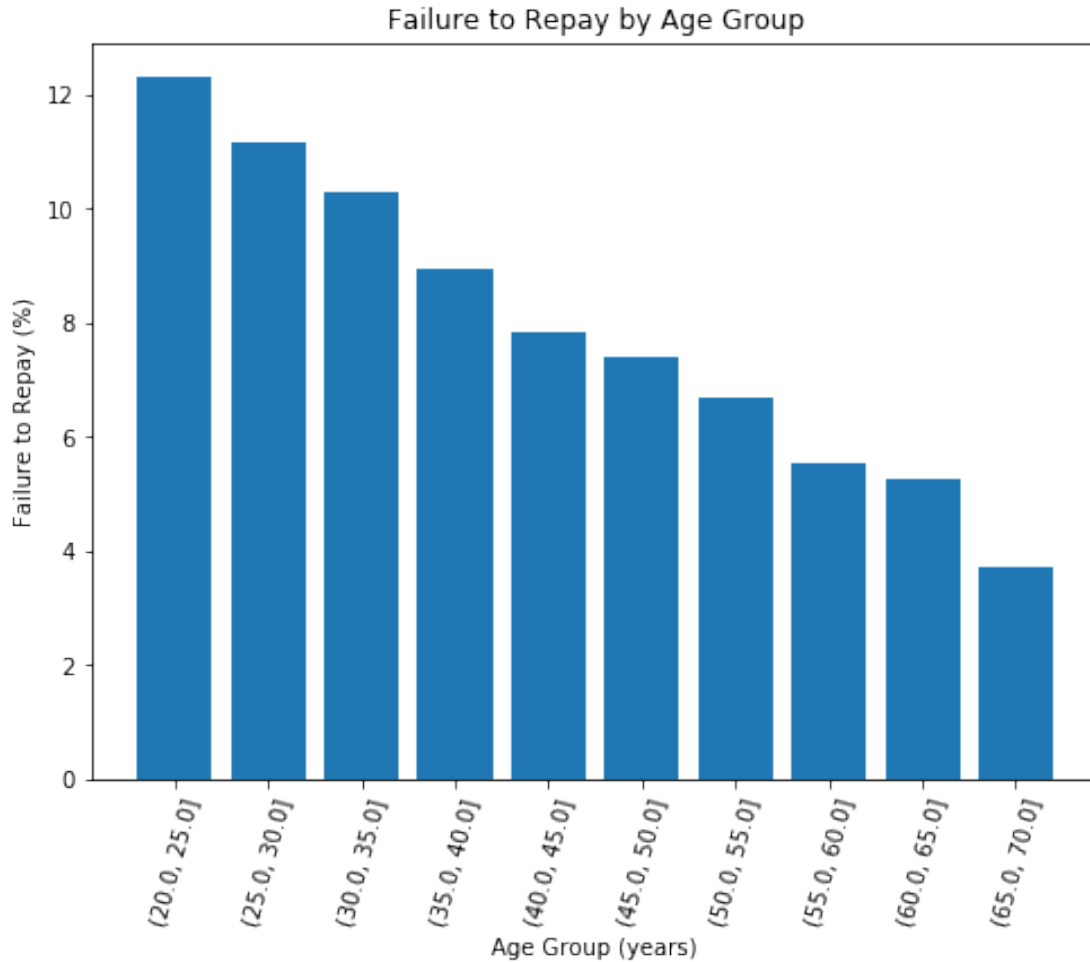
```
Out[40]:
```

	TARGET	DAYS_BIRTH	YEARS_BIRTH
YEARS_BINNED			
(20.0, 25.0]	0.123036	8532.795625	23.377522
(25.0, 30.0]	0.111436	10155.219250	27.822518
(30.0, 35.0]	0.102814	11854.848377	32.479037
(35.0, 40.0]	0.089414	13707.908253	37.555913
(40.0, 45.0]	0.078491	15497.661233	42.459346
(45.0, 50.0]	0.074171	17323.900441	47.462741
(50.0, 55.0]	0.066968	19196.494791	52.593136
(55.0, 60.0]	0.055314	20984.262742	57.491131
(60.0, 65.0]	0.052737	22780.547460	62.412459
(65.0, 70.0]	0.037270	24292.614340	66.555108

```
In [41]: plt.figure(figsize = (8, 6))

# Graph the age bins and the average of the target as a bar plot
plt.bar(age_groups.index.astype(str), 100 * age_groups['TARGET'])

# Plot labeling
plt.xticks(rotation = 75); plt.xlabel('Age Group (years)'); plt.ylabel('Failure to Repay')
plt.title('Failure to Repay by Age Group');
```



Younger applicants are more likely to not repay the loan.

4 Preparing data

4.1 Encoding Categorical Variables

A ML model can't deal with categorical features (except for some models such as LightGBM). One have to find a way to encode (represent) these variables as numbers. There are two main ways :

- Label encoding: assign each unique category in a categorical variable with an integer. No new columns are created. The problem with label encoding is that it gives the categories an arbitrary ordering.
- One-hot encoding: create a new column for each unique category in a categorical variable. Each observation receives a 1 in the column for its corresponding category and a 0 in all other new columns.

```
In [42]: app_train = pd.get_dummies(data=app_train, columns=obj_cols)
         app_test = pd.get_dummies(data=app_test, columns=obj_cols)
```

4.2 Aligning Training and Testing Data

Both the training and testing data should have the same features (columns). One-hot encoding can more columns in the one dataset because there were some categorical variables with categories not represented in the other dataset. In order to remove the columns in the training data that are not in the testing data, one need to align the dataframes.

```
In [43]: # back up of the target / need to keep this information
y = app_train.TARGET
app_train = app_train.drop(columns=['TARGET'])
```

```
In [44]: app_train, app_test = app_train.align(app_test, join = 'inner', axis = 1)
```

```
In [45]: app_train.shape, app_test.shape
```

```
Out[45]: ((307511, 168), (307511, 168))
```

4.3 Scaling values

```
In [46]: feat_to_scale = list(float_cols).copy()
feat_to_scale.extend(['CNT_CHILDREN', 'DAYS_BIRTH', 'DAYS_EMPLOYED', 'DAYS_ID_PUBLISH'])
feat_to_scale
```

```
Out[46]: ['AMT_INCOME_TOTAL',
'AMT_CREDIT',
'AMT_ANNUITY',
'AMT_GOODS_PRICE',
'REGION_POPULATION_RELATIVE',
'DAYS_REGISTRATION',
'CNT_FAM_MEMBERS',
'EXT_SOURCE_2',
'EXT_SOURCE_3',
'OBS_30_CNT_SOCIAL_CIRCLE',
'DEF_30_CNT_SOCIAL_CIRCLE',
'OBS_60_CNT_SOCIAL_CIRCLE',
'DEF_60_CNT_SOCIAL_CIRCLE',
'DAYS_LAST_PHONE_CHANGE',
'AMT_REQ_CREDIT_BUREAU_HOUR',
'AMT_REQ_CREDIT_BUREAU_DAY',
'AMT_REQ_CREDIT_BUREAU_WEEK',
'AMT_REQ_CREDIT_BUREAU_MON',
'AMT_REQ_CREDIT_BUREAU_QRT',
'AMT_REQ_CREDIT_BUREAU_YEAR',
'CNT_CHILDREN',
'DAYS_BIRTH',
'DAYS_EMPLOYED',
'DAYS_ID_PUBLISH',
'HOURL_APPR_PROCESS_START']
```

```

In [47]: scaler = StandardScaler()
         app_train[feat_to_scale] = scaler.fit_transform(app_train[feat_to_scale])
         app_test[feat_to_scale] = scaler.fit_transform(app_test[feat_to_scale])
         app_train.head()

/home/sunflowa/anaconda3/lib/python3.7/site-packages/sklearn/preprocessing/data.py:625: DataConversionWarning:
  return self.partial_fit(X, y)
/home/sunflowa/anaconda3/lib/python3.7/site-packages/sklearn/base.py:462: DataConversionWarning:
  return self.fit(X, **fit_params).transform(X)
/home/sunflowa/anaconda3/lib/python3.7/site-packages/sklearn/preprocessing/data.py:625: DataConversionWarning:
  return self.partial_fit(X, y)
/home/sunflowa/anaconda3/lib/python3.7/site-packages/sklearn/base.py:462: DataConversionWarning:
  return self.fit(X, **fit_params).transform(X)

```

```

Out [47]:
SK_ID_CURR  CNT_CHILDREN  AMT_INCOME_TOTAL  AMT_CREDIT  AMT_ANNUITY  \
0      100002      -0.577538         0.142129    -0.478095    -0.166143
1      100003      -0.577538         0.426792     1.725450     0.592683
2      100004      -0.577538        -0.427196    -1.152888    -1.404669
3      100006      -0.577538        -0.142533    -0.711430     0.177874
4      100007      -0.577538        -0.199466    -0.213734    -0.361749

AMT_GOODS_PRICE  REGION_POPULATION_RELATIVE  DAYS_BIRTH  DAYS_EMPLOYED  \
0      -0.507236         -0.149452    -1.506880     0.755835
1       1.600873        -1.252750     0.166821     0.497899
2      -1.092145        -0.783451     0.689509     0.948701
3      -0.653463        -0.928991     0.680114    -0.368597
4      -0.068554         0.563570     0.892535    -0.368129

DAYS_REGISTRATION  DAYS_ID_PUBLISH  FLAG_MOBIL  FLAG_EMP_PHONE  \
0         0.379837         0.579154         1         1
1         1.078697         1.790855         1         1
2         0.206116         0.306869         1         1
3        -1.375829         0.369143         1         1
4         0.191639        -0.307263         1         1

FLAG_WORK_PHONE  FLAG_CONT_MOBILE  FLAG_PHONE  FLAG_EMAIL  CNT_FAM_MEMBERS  \
0              0              1         1         0      -1.265722
1              0              1         1         0      -0.167638
2              1              1         1         0      -1.265722
3              0              1         0         0      -0.167638
4              0              1         0         0      -1.265722

REGION_RATING_CLIENT  REGION_RATING_CLIENT_W_CITY  HOUR_APPR_PROCESS_START  \
0              2              2      -0.631821
1              1              1      -0.325620
2              2              2     -0.938022
3              2              2      1.511587

```

4	2	2	-0.325620
---	---	---	-----------

	REG_REGION_NOT_LIVE_REGION	REG_REGION_NOT_WORK_REGION	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	LIVE_REGION_NOT_WORK_REGION	REG_CITY_NOT_LIVE_CITY	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	REG_CITY_NOT_WORK_CITY	LIVE_CITY_NOT_WORK_CITY	EXT_SOURCE_2	\
0	0	0	-1.317940	
1	0	0	0.564482	
2	0	0	0.216948	
3	0	0	0.712205	
4	1	1	-1.004691	

	EXT_SOURCE_3	OBS_30_CNT_SOCIAL_CIRCLE	DEF_30_CNT_SOCIAL_CIRCLE	\
0	-2.153651	0.242861	4.163504	
1	0.112063	-0.174085	-0.320480	
2	1.223975	-0.591031	-0.320480	
3	0.112063	0.242861	-0.320480	
4	0.112063	-0.591031	-0.320480	

	OBS_60_CNT_SOCIAL_CIRCLE	DEF_60_CNT_SOCIAL_CIRCLE	DAYS_LAST_PHONE_CHANGE	\
0	0.252132	5.253260	-0.206992	
1	-0.168527	-0.275663	0.163107	
2	-0.589187	-0.275663	0.178831	
3	0.252132	-0.275663	0.418306	
4	-0.589187	-0.275663	-0.173126	

	FLAG_DOCUMENT_2	FLAG_DOCUMENT_3	FLAG_DOCUMENT_4	FLAG_DOCUMENT_5	\
0	0	1	0	0	
1	0	1	0	0	
2	0	0	0	0	
3	0	1	0	0	
4	0	0	0	0	

	FLAG_DOCUMENT_6	FLAG_DOCUMENT_7	FLAG_DOCUMENT_8	FLAG_DOCUMENT_9	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	

3	0	0	0	0
4	0	0	1	0

	FLAG_DOCUMENT_10	FLAG_DOCUMENT_11	FLAG_DOCUMENT_12	FLAG_DOCUMENT_13	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	

	FLAG_DOCUMENT_14	FLAG_DOCUMENT_15	FLAG_DOCUMENT_16	FLAG_DOCUMENT_17	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	

	FLAG_DOCUMENT_18	FLAG_DOCUMENT_19	FLAG_DOCUMENT_20	FLAG_DOCUMENT_21	\
0	0	0	0	0	
1	0	0	0	0	
2	0	0	0	0	
3	0	0	0	0	
4	0	0	0	0	

	AMT_REQ_CREDIT_BUREAU_HOUR	AMT_REQ_CREDIT_BUREAU_DAY	\
0	-0.070987	-0.058766	
1	-0.070987	-0.058766	
2	-0.070987	-0.058766	
3	-0.070987	-0.058766	
4	-0.070987	-0.058766	

	AMT_REQ_CREDIT_BUREAU_WEEK	AMT_REQ_CREDIT_BUREAU_MON	\
0	-0.155837	-0.269947	
1	-0.155837	-0.269947	
2	-0.155837	-0.269947	
3	-0.155837	-0.269947	
4	-0.155837	-0.269947	

	AMT_REQ_CREDIT_BUREAU_QRT	AMT_REQ_CREDIT_BUREAU_YEAR	DAYS_EMPLOYED_ANOM	\
0	-0.30862	-0.440926	False	
1	-0.30862	-1.007331	False	
2	-0.30862	-1.007331	False	
3	-0.30862	-0.440926	False	
4	-0.30862	-1.007331	False	

	NAME_CONTRACT_TYPE_Cash loans	NAME_CONTRACT_TYPE_Revolving loans	\
0	1	0	
1	1	0	

2	0	1
3	1	0
4	1	0

	CODE_GENDER_F	CODE_GENDER_M	CODE_GENDER_XNA	FLAG_OWN_CAR_N	\
0	0	1	0	1	
1	1	0	0	1	
2	0	1	0	0	
3	1	0	0	1	
4	0	1	0	1	

	FLAG_OWN_CAR_Y	FLAG_OWN_REALTY_N	FLAG_OWN_REALTY_Y	\
0	0	0	1	
1	0	1	0	
2	1	0	1	
3	0	0	1	
4	0	0	1	

	NAME_TYPE_SUITE_Children	NAME_TYPE_SUITE_Family	\
0	0	0	
1	0	1	
2	0	0	
3	0	0	
4	0	0	

	NAME_TYPE_SUITE_Group of people	NAME_TYPE_SUITE_Not specified	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	NAME_TYPE_SUITE_Other_A	NAME_TYPE_SUITE_Other_B	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	NAME_TYPE_SUITE_Spouse, partner	NAME_TYPE_SUITE_Unaccompanied	\
0	0	1	
1	0	0	
2	0	1	
3	0	1	
4	0	1	

	NAME_INCOME_TYPE_Businessman	NAME_INCOME_TYPE_Commercial associate	\
0	0	0	

1	0	0
2	0	0
3	0	0
4	0	0

	NAME_INCOME_TYPE_Maternity leave	NAME_INCOME_TYPE_Pensioner \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

	NAME_INCOME_TYPE_State servant	NAME_INCOME_TYPE_Student \
0	0	0
1	1	0
2	0	0
3	0	0
4	0	0

	NAME_INCOME_TYPE_Unemployed	NAME_INCOME_TYPE_Working \
0	0	1
1	0	0
2	0	1
3	0	1
4	0	1

	NAME_EDUCATION_TYPE_Academic degree	NAME_EDUCATION_TYPE_Higher education \
0	0	0
1	0	1
2	0	0
3	0	0
4	0	0

	NAME_EDUCATION_TYPE_Incomplete higher	NAME_EDUCATION_TYPE_Lower secondary \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

	NAME_EDUCATION_TYPE_Secondary / secondary special \
0	1
1	0
2	1
3	1
4	1

	NAME_FAMILY_STATUS_Civil marriage	NAME_FAMILY_STATUS_Married \
--	-----------------------------------	------------------------------

0	0	0
1	0	1
2	0	0
3	1	0
4	0	0

	NAME_FAMILY_STATUS_Separated	NAME_FAMILY_STATUS_Single / not married \
0	0	1
1	0	0
2	0	1
3	0	0
4	0	1

	NAME_FAMILY_STATUS_Unknown	NAME_FAMILY_STATUS_Widow \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

	NAME_HOUSING_TYPE_Co-op apartment	NAME_HOUSING_TYPE_House / apartment \
0	0	1
1	0	1
2	0	1
3	0	1
4	0	1

	NAME_HOUSING_TYPE_Municipal apartment	NAME_HOUSING_TYPE_Office apartment \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

	NAME_HOUSING_TYPE_Rented apartment	NAME_HOUSING_TYPE_With parents \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

	WEEKDAY_APPR_PROCESS_START_FRIDAY	WEEKDAY_APPR_PROCESS_START_MONDAY \
0	0	0
1	0	1
2	0	1
3	0	0
4	0	0

	WEEKDAY_APPR_PROCESS_START_SATURDAY	WEEKDAY_APPR_PROCESS_START_SUNDAY	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	WEEKDAY_APPR_PROCESS_START_THURSDAY	WEEKDAY_APPR_PROCESS_START_TUESDAY	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	1	0	

	WEEKDAY_APPR_PROCESS_START_WEDNESDAY	ORGANIZATION_TYPE_Advertising	\
0	1	0	
1	0	0	
2	0	0	
3	1	0	
4	0	0	

	ORGANIZATION_TYPE_Agriculture	ORGANIZATION_TYPE_Bank	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	ORGANIZATION_TYPE_Business Entity Type 1	\
0	0	
1	0	
2	0	
3	0	
4	0	

	ORGANIZATION_TYPE_Business Entity Type 2	\
0	0	
1	0	
2	0	
3	0	
4	0	

	ORGANIZATION_TYPE_Business Entity Type 3	ORGANIZATION_TYPE_Cleaning	\
0	1	0	
1	0	0	
2	0	0	
3	1	0	
4	0	0	

	ORGANIZATION_TYPE_Construction	ORGANIZATION_TYPE_Culture \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

	ORGANIZATION_TYPE_Electricity	ORGANIZATION_TYPE_Emergency \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

	ORGANIZATION_TYPE_Government	ORGANIZATION_TYPE_Hotel \
0	0	0
1	0	0
2	1	0
3	0	0
4	0	0

	ORGANIZATION_TYPE_Housing	ORGANIZATION_TYPE_Industry: type 1 \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

	ORGANIZATION_TYPE_Industry: type 10	ORGANIZATION_TYPE_Industry: type 11 \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

	ORGANIZATION_TYPE_Industry: type 12	ORGANIZATION_TYPE_Industry: type 13 \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

	ORGANIZATION_TYPE_Industry: type 2	ORGANIZATION_TYPE_Industry: type 3 \
0	0	0
1	0	0
2	0	0
3	0	0

4	0	0
ORGANIZATION_TYPE_Industry: type 4	ORGANIZATION_TYPE_Industry: type 5	\
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
ORGANIZATION_TYPE_Industry: type 6	ORGANIZATION_TYPE_Industry: type 7	\
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
ORGANIZATION_TYPE_Industry: type 8	ORGANIZATION_TYPE_Industry: type 9	\
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
ORGANIZATION_TYPE_Insurance	ORGANIZATION_TYPE_Kindergarten	\
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
ORGANIZATION_TYPE_Legal Services	ORGANIZATION_TYPE_Medicine	\
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
ORGANIZATION_TYPE_Military	ORGANIZATION_TYPE_Mobile	\
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
ORGANIZATION_TYPE_Other	ORGANIZATION_TYPE_Police	\
0	0	0
1	0	0
2	0	0

3	0	0
4	0	0

	ORGANIZATION_TYPE_Postal	ORGANIZATION_TYPE_Realtor	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	ORGANIZATION_TYPE_Religion	ORGANIZATION_TYPE_Restaurant	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	1	0	

	ORGANIZATION_TYPE_School	ORGANIZATION_TYPE_Security	\
0	0	0	
1	1	0	
2	0	0	
3	0	0	
4	0	0	

	ORGANIZATION_TYPE_Security Ministries	ORGANIZATION_TYPE_Self-employed	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	ORGANIZATION_TYPE_Services	ORGANIZATION_TYPE_Telecom	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	ORGANIZATION_TYPE_Trade: type 1	ORGANIZATION_TYPE_Trade: type 2	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	ORGANIZATION_TYPE_Trade: type 3	ORGANIZATION_TYPE_Trade: type 4	\
0	0	0	
1	0	0	

2	0	0
3	0	0
4	0	0

ORGANIZATION_TYPE_Trade: type 5		ORGANIZATION_TYPE_Trade: type 6 \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

ORGANIZATION_TYPE_Trade: type 7		ORGANIZATION_TYPE_Transport: type 1 \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

ORGANIZATION_TYPE_Transport: type 2		ORGANIZATION_TYPE_Transport: type 3 \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

ORGANIZATION_TYPE_Transport: type 4		ORGANIZATION_TYPE_University \
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0

ORGANIZATION_TYPE_XNA	
0	0
1	0
2	0
3	0
4	0

4.4 Splitting training / test datasets

from app_train in order to make few predictions before submission & select models

In [48]: `X_train, X_test, y_train, y_test = train_test_split(app_train, y, test_size=0.2)`

5 Base line

5.1 Metric: ROC AUC

more infos on [the Receiver Operating Characteristic Area Under the Curve \(ROC AUC, also sometimes called AUROC\)](#).

The [Reciever Operating Characteristic \(ROC\) curve](#) graphs the true positive rate versus the false positive rate:

A single line on the graph indicates the curve for a single model, and movement along a line indicates changing the threshold used for classifying a positive instance. The threshold starts at 0 in the upper right to and goes to 1 in the lower left. A curve that is to the left and above another curve indicates a better model. For example, the blue model is better than the red one (which is better than the black diagonal line which indicates a naive random guessing model).

The Area Under the Curve (AUC) is the integral of the curve. This metric is between 0 and 1 with a better model scoring higher. A model that simply guesses at random will have an ROC AUC of 0.5.

When we measure a classifier according to the ROC AUC, we do not generate 0 or 1 predictions, but rather a probability between 0 and 1.

When we get into problems with inbalanced classes, accuracy is not the best metric. A model with a high ROC AUC will also have a high accuracy, but the ROC AUC is a better representation of model performance.

5.2 Random forrest

```
In [49]: # a simple RandomForest Classifier without CV
         rf = RandomForestClassifier(n_estimators=50)
         rf.fit(X_train, y_train)
         y_pred = rf.predict(X_test)
         roc_auc_score(y_test, y_pred)
```

```
Out[49]: 0.5011674887648608
```

The predictions must be in the format shown in the sample_submission.csv file, where there are only two columns: SK_ID_CURR and TARGET. Let's create a dataframe in this format from the test set and the predictions called submit.

```
In [50]: def submit(model, csv_name):

         # fit on the whole dataset of train
         model.fit(app_train, y)

         # Make predictions & make sure to select the second column only
         result = model.predict_proba(app_test)[:, 1]

         submit = app_test[['SK_ID_CURR']]
         submit['TARGET'] = result

         # Save the submission to a csv file
         submit.to_csv(csv_name, index = False)
```

```
In [51]: # submit(rf, 'random_forrest_clf.csv')
```

The random forrest model should score around 0.58329 when submitted which is not really good, because just above 0.5 i.e a random classifier...

5.3 Feature Importances

```
In [52]: importances = rf.feature_importances_  
std = np.std([tree.feature_importances_ for tree in rf.estimators_], axis=0)  
indices = np.argsort(importances)[::-1]  
  
# Print the feature ranking  
print("Feature ranking:")  
  
for f in range(app_train.shape[1]):  
    print("%d. feature %d (%f)" % (f + 1, indices[f], importances[indices[f]]))
```

Feature ranking:

```
1. feature 27 (0.064841)  
2. feature 28 (0.059712)  
3. feature 7 (0.046094)  
4. feature 10 (0.045429)  
5. feature 9 (0.045176)  
6. feature 0 (0.044173)  
7. feature 4 (0.041803)  
8. feature 8 (0.041481)  
9. feature 33 (0.041030)  
10. feature 3 (0.039948)  
11. feature 2 (0.035454)  
12. feature 6 (0.035247)  
13. feature 5 (0.033751)  
14. feature 20 (0.031268)  
15. feature 59 (0.022036)  
16. feature 31 (0.017715)  
17. feature 29 (0.017527)  
18. feature 17 (0.013756)  
19. feature 1 (0.010035)  
20. feature 58 (0.008328)  
21. feature 57 (0.007471)  
22. feature 18 (0.006926)  
23. feature 69 (0.006884)  
24. feature 19 (0.006881)  
25. feature 15 (0.006740)  
26. feature 92 (0.006712)  
27. feature 68 (0.006608)  
28. feature 30 (0.006549)  
29. feature 115 (0.006527)  
30. feature 108 (0.006369)
```

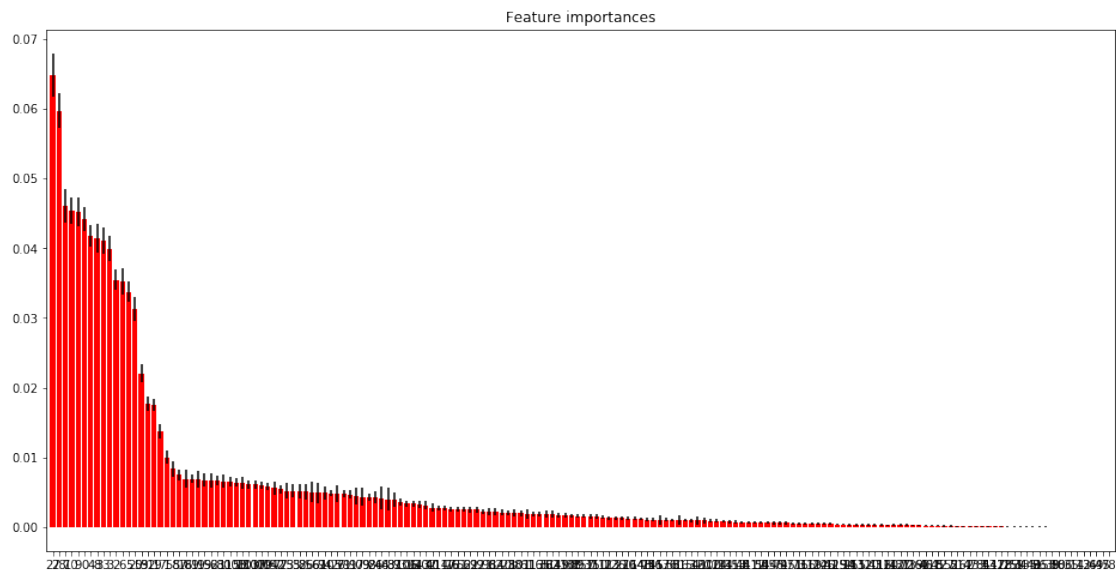
31. feature 13 (0.006358)
32. feature 103 (0.006115)
33. feature 107 (0.006110)
34. feature 109 (0.006067)
35. feature 104 (0.005811)
36. feature 152 (0.005601)
37. feature 77 (0.005453)
38. feature 25 (0.005225)
39. feature 35 (0.005202)
40. feature 32 (0.005143)
41. feature 85 (0.005102)
42. feature 66 (0.004996)
43. feature 67 (0.004896)
44. feature 94 (0.004894)
45. feature 105 (0.004844)
46. feature 26 (0.004821)
47. feature 71 (0.004800)
48. feature 91 (0.004709)
49. feature 90 (0.004508)
50. feature 79 (0.004343)
51. feature 98 (0.004327)
52. feature 24 (0.004303)
53. feature 64 (0.004122)
54. feature 63 (0.004009)
55. feature 87 (0.003993)
56. feature 93 (0.003624)
57. feature 106 (0.003384)
58. feature 16 (0.003367)
59. feature 143 (0.003299)
60. feature 102 (0.003096)
61. feature 40 (0.002746)
62. feature 114 (0.002712)
63. feature 117 (0.002707)
64. feature 76 (0.002608)
65. feature 161 (0.002586)
66. feature 56 (0.002539)
67. feature 22 (0.002522)
68. feature 99 (0.002518)
69. feature 121 (0.002230)
70. feature 96 (0.002213)
71. feature 82 (0.002195)
72. feature 140 (0.002136)
73. feature 23 (0.002091)
74. feature 88 (0.002060)
75. feature 101 (0.001985)
76. feature 61 (0.001886)
77. feature 113 (0.001883)
78. feature 165 (0.001881)

79. feature 38 (0.001878)
80. feature 62 (0.001859)
81. feature 149 (0.001758)
82. feature 130 (0.001690)
83. feature 138 (0.001645)
84. feature 89 (0.001614)
85. feature 157 (0.001587)
86. feature 37 (0.001502)
87. feature 150 (0.001492)
88. feature 111 (0.001400)
89. feature 123 (0.001298)
90. feature 126 (0.001292)
91. feature 21 (0.001287)
92. feature 70 (0.001248)
93. feature 164 (0.001238)
94. feature 148 (0.001203)
95. feature 75 (0.001104)
96. feature 145 (0.001089)
97. feature 167 (0.001063)
98. feature 136 (0.001062)
99. feature 55 (0.001006)
100. feature 81 (0.001001)
101. feature 163 (0.000992)
102. feature 54 (0.000963)
103. feature 12 (0.000949)
104. feature 60 (0.000938)
105. feature 100 (0.000927)
106. feature 124 (0.000895)
107. feature 134 (0.000826)
108. feature 131 (0.000789)
109. feature 153 (0.000724)
110. feature 48 (0.000719)
111. feature 141 (0.000694)
112. feature 112 (0.000693)
113. feature 50 (0.000687)
114. feature 144 (0.000681)
115. feature 156 (0.000631)
116. feature 74 (0.000601)
117. feature 97 (0.000592)
118. feature 151 (0.000589)
119. feature 73 (0.000482)
120. feature 166 (0.000461)
121. feature 119 (0.000460)
122. feature 122 (0.000450)
123. feature 146 (0.000447)
124. feature 41 (0.000426)
125. feature 129 (0.000396)
126. feature 154 (0.000395)

```
127. feature 14 (0.000376)
128. feature 155 (0.000343)
129. feature 132 (0.000341)
130. feature 110 (0.000311)
131. feature 43 (0.000310)
132. feature 116 (0.000301)
133. feature 120 (0.000282)
134. feature 142 (0.000280)
135. feature 137 (0.000279)
136. feature 72 (0.000267)
137. feature 139 (0.000261)
138. feature 160 (0.000257)
139. feature 46 (0.000214)
140. feature 118 (0.000208)
141. feature 45 (0.000186)
142. feature 127 (0.000184)
143. feature 52 (0.000161)
144. feature 51 (0.000144)
145. feature 162 (0.000106)
146. feature 47 (0.000101)
147. feature 133 (0.000101)
148. feature 84 (0.000093)
149. feature 53 (0.000089)
150. feature 147 (0.000088)
151. feature 128 (0.000085)
152. feature 125 (0.000069)
153. feature 159 (0.000063)
154. feature 34 (0.000057)
155. feature 135 (0.000053)
156. feature 49 (0.000051)
157. feature 86 (0.000044)
158. feature 158 (0.000037)
159. feature 39 (0.000020)
160. feature 80 (0.000020)
161. feature 83 (0.000001)
162. feature 65 (0.000001)
163. feature 11 (0.000000)
164. feature 42 (0.000000)
165. feature 36 (0.000000)
166. feature 44 (0.000000)
167. feature 95 (0.000000)
168. feature 78 (0.000000)
```


```
In [53]: # Plot the feature importances of the rf
plt.figure(figsize=(16, 8))
plt.title("Feature importances")
plt.bar(range(app_train.shape[1]), importances[indices], color="r", yerr=std[indices])
```

```
plt.xticks(range(app_train.shape[1]), indices)
plt.xlim([-1, app_train.shape[1]])
plt.show()
```



```
In [54]: (pd.Series(rf.feature_importances_, index=app_train.columns)
          .nlargest(15)
          .plot(kind='barh'))
```

```
Out[54]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7b8d800be0>
```



5.4 Random forrest with a cross validation

```
In [55]: rf_cv = RandomForestClassifier()
         scores = cross_val_score(rf_cv, X_train, y_train, cv=5, scoring='roc_auc', n_jobs=-1)
         scores

Out[55]: array([0.63219173, 0.63231896, 0.62319801, 0.62533242, 0.62668538])

In [57]: rf_cv.fit(X_train, y_train)
         roc_auc_score(y_test, rf_cv.predict(X_test))

Out[57]: 0.5036931060869384

In [58]: #!pip install kaggle

In [59]: #!kaggle competitions submit -c home-credit-default-risk -f randomforest_baseline.csv
```

6 More advanced models

6.1 LightGBM

```
In [62]: lgbm = lgb.LGBMClassifier(random_state = 50)
         lgbm.fit(X_train, y_train, eval_metric = 'auc')
         roc_auc_score(y_train, lgbm.predict(X_train))
```

Out[62]: 0.5108945200660795

```
In [63]: roc_auc_score(y_test, lgbm.predict(X_test))
```

Out[63]: 0.5069964776833696

Different tests on hyperparameters and results:

- underfitting / high bias -> let's try to complicate the model
- max_depth = 7/11 or objective = 'binary' -> scores 0.508 / 0.508
- n_estimators=1000 -> scores 0.57 / 0.511
- class_weight = 'balanced' -> scores 0.71 / 0.68
- reg_alpha = 0.1, reg_lambda = 0.1 -> no influence

```
In [64]: lgbm = lgb.LGBMClassifier(random_state = 50, n_jobs = -1, class_weight = 'balanced')
         lgbm.fit(X_train, y_train, eval_metric = 'auc')
         roc_auc_score(y_train, lgbm.predict(X_train))
```

Out[64]: 0.7121220817520526

```
In [65]: roc_auc_score(y_test, lgbm.predict(X_test))
```

Out[65]: 0.6846561563080866


```
In [66]: def submit_func(model, X_Test, file_name):
        model.fit(app_train, y)
        result = model.predict_proba(app_test)[: , 1]
        submit = app_test[['SK_ID_CURR']]
        submit['TARGET'] = result
        print(submit.head())
        print(submit.shape)
        submit.to_csv(file_name + '.csv', index=False)
```

```
In [67]: submit_func(lgbm, app_test, 'lgbm_submission')
```

/home/sunflowa/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:5: SettingWithCopyWarning: A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>
"""

	SK_ID_CURR	TARGET
0	100002	0.876184
1	100003	0.237339
2	100004	0.293612
3	100006	0.427514
4	100007	0.634639

(307511, 2)

submission -> 0.72057

7 Using XGBoost and weighted classes

As said earlier, there are far more 0 than 1 in the target column. This is an [imbalanced class problem].(<http://www.chioka.in/class-imbalance-problem/>).

It's a common problem affecting ML due to having disproportionate number of class instances in practice. This is why the ROC AUC metric suits our needs here. There are 2 class of approaches out there to deal with this problem:

- 1) sampling based, that can be broken into three major categories:
 - a) over sampling
 - b) under sampling
 - c) hybrid of oversampling and undersampling.
- 2) cost function based.

With default or few changes in hyperparameters

- base score : 0.50 / 0.709
- max_delta_step=2 -> unchanged
- with ratio : 0.68 / 0.71

```
In [68]: y.shape[0], y.sum()
```

```
Out[68]: (307511, 24825)
```

```
In [69]: ratio = (y.shape[0] - y.sum()) / y.sum()
         ratio
```

```
Out[69]: 11.387150050352467
```

```
In [70]: xgb_model = xgb.XGBClassifier(objective="binary:logistic", random_state=50, eval_metric='auc',
                                     max_delta_step=2, scale_pos_weight=20)
         xgb_model.fit(X_train, y_train)
         roc_auc_score(y_train, xgb_model.predict(X_train))
```

```
Out[70]: 0.656946528564419
```

```
In [71]: roc_auc_score(y_test, xgb_model.predict(X_test))
```

```
Out[71]: 0.6488130660302404
```

For common cases when the dataset is extremely imbalanced, this can affect the training of XGBoost model, and there are two ways to improve it.

If you care only about the overall performance metric (AUC) of your prediction Balance the positive and negative weights via scale_pos_weight Use AUC for evaluation

If you care about predicting the right probability In such a case, you cannot re-balance the dataset Set parameter max_delta_step to a finite number (say 1) to help convergence

```
In [72]: submit_func(xgb_model, app_test, 'xgb_submission')
```

```
/home/sunflowa/anaconda3/lib/python3.7/site-packages/ipykernel_launcher.py:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

```
SK_ID_CURR  TARGET
0      100002  0.908168
1      100003  0.408862
2      100004  0.452180
3      100006  0.607726
4      100007  0.763639
(307511, 2)
```

submission -> 0.72340

8 Credits / side notes

[Will Koehrsen](#) for many interesting tips in his kernel !

This notebook is intended to be an introduction to machine learning. So many things are missing or can be done better, such as :

- Using function to clean / prepare the data
- Exploring the other tables and select other columns that can be relevant
- Doing more feature engineering, this will lead to a better score !