

# California Housing Prices

June 27, 2019

## 1 California Housing Prices

Median house prices for California districts derived from the 1990 census.

Photo by [Vita Vilcina](#)

### 1.1 Context

This is the dataset used in the second chapter of Aurélien Géron's recent book 'Hands-On Machine learning with Scikit-Learn and TensorFlow'. It serves as an excellent introduction to implementing machine learning algorithms because it requires rudimentary data cleaning, has an easily understandable list of variables and sits at an optimal size between being too toyish and too cumbersome.

The data contains information from the 1990 California census. So although it may not help you with predicting current housing prices like the Zillow Zestimate dataset, it does provide an accessible introductory dataset for teaching people about the basics of machine learning.

### 1.2 Acknowledgements

Please refer to the [Kaggle challenge web page](#)

### 1.3 Inspiration

predict a real estate price

---

## 2 Exploratory Data Analysis

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import os
```

```
In [2]: import folium
```

```
In [3]: from sklearn.model_selection import train_test_split
        from sklearn.metrics import mean_squared_error
        from sklearn.linear_model import Lasso, LinearRegression, Ridge, RANSACRegressor, SGDRegressor
        from sklearn.ensemble import AdaBoostRegressor
        from sklearn.svm import SVR
```

```
In [4]: file_path = os.path.join('input', 'house_big.csv')
        df = pd.read_csv(file_path)
        df.head()
```

```
Out[4]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	\
0	-122.23	37.88	41.0	880.0	129.0	
1	-122.22	37.86	21.0	7099.0	1106.0	
2	-122.24	37.85	52.0	1467.0	190.0	
3	-122.25	37.85	52.0	1274.0	235.0	
4	-122.25	37.85	52.0	1627.0	280.0	

	population	households	median_income	median_house_value	ocean_proximity
0	322.0	126.0	8.3252	452600.0	NEAR BAY
1	2401.0	1138.0	8.3014	358500.0	NEAR BAY
2	496.0	177.0	7.2574	352100.0	NEAR BAY
3	558.0	219.0	5.6431	341300.0	NEAR BAY
4	565.0	259.0	3.8462	342200.0	NEAR BAY

```
In [5]: df.shape
```

```
Out[5]: (20640, 10)
```

## 2.1 Content

The data pertains to the houses found in a given California district and some summary stats about them based on the 1990 census data. Be warned the data aren't cleaned so there are some preprocessing steps required! The columns are as follows, their names are pretty self explanatory: \* longitude \* latitude \* housing\_median\_age \* total\_rooms \* total\_bedrooms \* population \* households \* median\_income \* median\_house\_value \* ocean\_proximity

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
longitude          20640 non-null float64
latitude           20640 non-null float64
housing_median_age 20640 non-null float64
total_rooms         20640 non-null float64
total_bedrooms      20433 non-null float64
population          20640 non-null float64
households          20640 non-null float64
median_income       20640 non-null float64
```

```

median_house_value    20640 non-null float64
ocean_proximity       20640 non-null object
dtypes: float64(9), object(1)
memory usage: 1.6+ MB

```

There are few missing value int the 'total\_bedrooms' column. Now let's see the basic stats for the numerical columns:

```
In [7]: df.describe()
```

```

Out [7]:
      longitude  latitude  housing_median_age  total_rooms  \
count  20640.000000  20640.000000      20640.000000  20640.000000
mean   -119.569704    35.631861         28.639486   2635.763081
std      2.003532     2.135952         12.585558   2181.615252
min    -124.350000    32.540000          1.000000     2.000000
25%    -121.800000    33.930000         18.000000   1447.750000
50%    -118.490000    34.260000         29.000000   2127.000000
75%    -118.010000    37.710000         37.000000   3148.000000
max     -114.310000    41.950000         52.000000  39320.000000

      total_bedrooms  population  households  median_income  \
count  20433.000000  20640.000000  20640.000000  20640.000000
mean     537.870553   1425.476744    499.539680     3.870671
std     421.385070   1132.462122    382.329753     1.899822
min       1.000000     3.000000     1.000000     0.499900
25%     296.000000    787.000000    280.000000     2.563400
50%     435.000000   1166.000000    409.000000     3.534800
75%     647.000000   1725.000000    605.000000     4.743250
max    6445.000000  35682.000000   6082.000000    15.000100

      median_house_value
count      20640.000000
mean     206855.816909
std     115395.615874
min       14999.000000
25%     119600.000000
50%     179700.000000
75%     264725.000000
max     500001.000000

```

```
In [8]: df.ocean_proximity.value_counts()
```

```

Out [8]: <1H OCEAN      9136
INLAND      6551
NEAR OCEAN   2658
NEAR BAY     2290
ISLAND         5
Name: ocean_proximity, dtype: int64

```

## 2.2 Cleaning data

```
In [9]: df.duplicated().sum()
```

```
Out[9]: 0
```

```
In [10]: df.isnull().sum()
```

```
Out[10]: longitude          0
         latitude           0
         housing_median_age  0
         total_rooms         0
         total_bedrooms     207
         population         0
         households         0
         median_income       0
         median_house_value  0
         ocean_proximity     0
         dtype: int64
```

```
In [11]: print(f'percentage of missing values: {df.total_bedrooms.isnull().sum() / df.shape[0]}
```

```
percentage of missing values: 1.00%
```

```
In [12]: df = df.fillna(df.median())
         df.isnull().sum()
```


```
Out[12]: longitude          0
         latitude           0
         housing_median_age  0
         total_rooms         0
         total_bedrooms     0
         population         0
         households         0
         median_income       0
         median_house_value  0
         ocean_proximity     0
         dtype: int64
```

## 2.3 Dealing with geospatial infos

Visualization of the data in a scatter plot in a "geographic way"


```
In [13]: sns.scatterplot(df.longitude, df.latitude)
```

```
Out[13]: <matplotlib.axes._subplots.AxesSubplot at 0x7f244cbeeb00>
```



Same plot but this time with a varying size of the data points based on population variable and a different color depending of the real estate price (median\_house\_value)

```
In [14]: sns.relplot(x="longitude", y="latitude", hue="median_house_value", size="population",  
                    sizes=(50, 700), data=df, height=8)  
plt.show()
```



```
In [15]: # Create a map with folium centered at the mean latitude and longitude
cali_map = folium.Map(location=[35.6, -117], zoom_start=6)
```

```
# Display the map
display(cali_map)
```

```
<folium.folium.Map at 0x7f244c886d68>
```

```
In [16]: # Add markers for each rows
for i in range(df.shape[0]):
    folium.Marker((float(df.iloc[i, 1]), float(df.iloc[i, 0]))).add_to(cali_map)
```

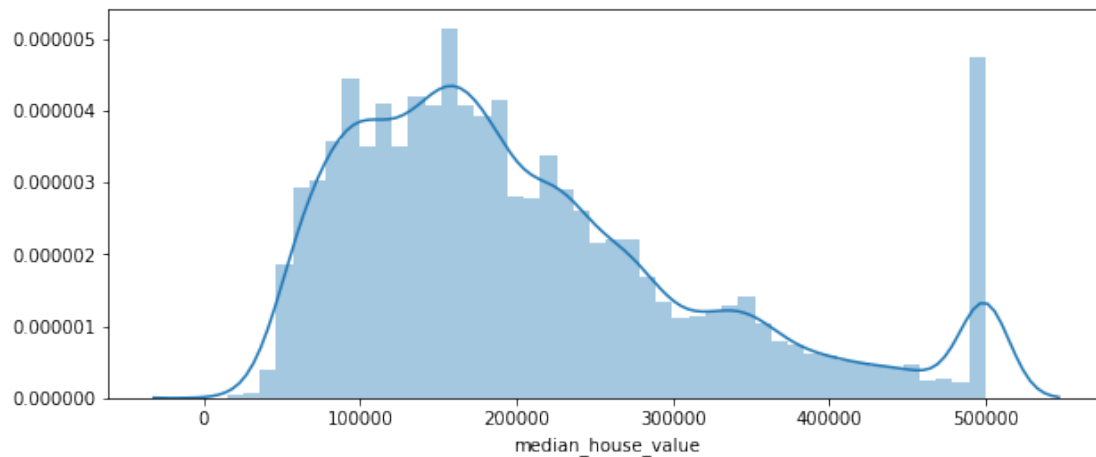
```
# Display the map
display(cali_map)
```

```
<folium.folium.Map at 0x7f244c886d68>
```

## 2.4 Target analysis

```
In [17]: plt.figure(figsize=(10, 4))
         sns.distplot(df.median_house_value)
         plt.show()
```

```
/home/sunflow/anaconda3/lib/python3.7/site-packages/scipy/stats/stats.py:1713: FutureWarning:
  return np.add.reduce(sorted[indexer] * weights, axis=axis) / sumval
```




Variations depending on the proximity with ocean

```
In [18]: df.ocean_proximity.unique()
```

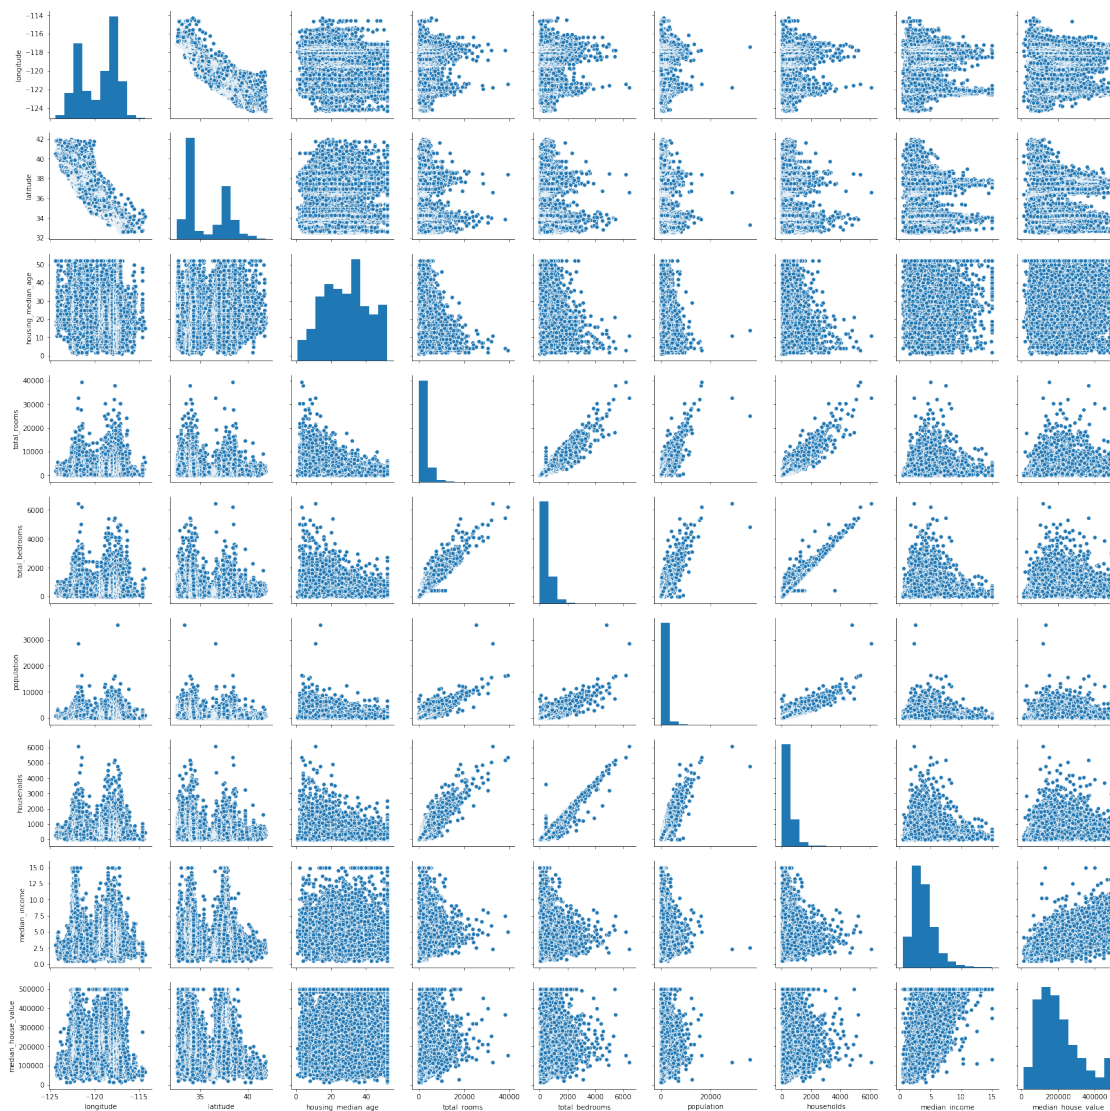
```
Out[18]: array(['NEAR BAY', '<1H OCEAN', 'INLAND', 'NEAR OCEAN', 'ISLAND'],
              dtype=object)
```

```
In [19]: plt.figure(figsize=(10, 4))
         for prox in df.ocean_proximity.unique():
             sns.kdeplot(data=df[df.ocean_proximity == prox].median_house_value)
         plt.legend(prox)
         plt.show()
```




## 2.5 Other analysis

```
In [20]: sns.pairplot(df)
plt.show()
```



```
In [21]: df.hist(figsize=(8, 8))
plt.show()
```





## 2.6 Correlations

```
In [22]: corr = df.corr()
         corr
```

```
Out [22]:
```

	longitude	latitude	housing_median_age	total_rooms	\
longitude	1.000000	-0.924664	-0.108197	0.044568	
latitude	-0.924664	1.000000	0.011173	-0.036100	
housing_median_age	-0.108197	0.011173	1.000000	-0.361262	
total_rooms	0.044568	-0.036100	-0.361262	1.000000	
total_bedrooms	0.069120	-0.066484	-0.319026	0.927058	
population	0.099773	-0.108785	-0.296244	0.857126	
households	0.055310	-0.071035	-0.302916	0.918484	
median_income	-0.015176	-0.079809	-0.119034	0.198050	
median_house_value	-0.045967	-0.144160	0.105623	0.134153	

	total_bedrooms	population	households	median_income	\
longitude	0.069120	0.099773	0.055310	-0.015176	
latitude	-0.066484	-0.108785	-0.071035	-0.079809	
housing_median_age	-0.319026	-0.296244	-0.302916	-0.119034	
total_rooms	0.927058	0.857126	0.918484	0.198050	
total_bedrooms	1.000000	0.873535	0.974366	-0.007617	
population	0.873535	1.000000	0.907222	0.004834	
households	0.974366	0.907222	1.000000	0.013033	
median_income	-0.007617	0.004834	0.013033	1.000000	
median_house_value	0.049457	-0.024650	0.065843	0.688075	

	median_house_value
longitude	-0.045967
latitude	-0.144160
housing_median_age	0.105623
total_rooms	0.134153
total_bedrooms	0.049457
population	-0.024650
households	0.065843
median_income	0.688075
median_house_value	1.000000


```
In [23]: # Generate a mask for the upper triangle
mask = np.zeros_like(corr, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True

# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(8, 6))

# Generate a custom diverging colormap
cmap = sns.diverging_palette(220, 10, as_cmap=True)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, vmax=.3, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5}, annot=True)

Out[23]: <matplotlib.axes._subplots.AxesSubplot at 0x7f243ec37898>
```



- lat and log are highly positively correlated
- total\_bedrooms, population and households are highly positively correlated too
- median\_income and median\_house\_value are also positively correlated

which make sense.

### 3 Models training and predictions

#### 3.1 Data preparation

Label encoding of categorical feature (ocean proximity)

```
In [24]: df = pd.get_dummies(data=df, columns=['ocean_proximity'], drop_first=False)
df.head()
```

```

Out [24]:  longitude  latitude  housing_median_age  total_rooms  total_bedrooms  \
0      -122.23      37.88                41.0          880.0          129.0
1      -122.22      37.86                21.0         7099.0         1106.0
2      -122.24      37.85                52.0         1467.0          190.0
3      -122.25      37.85                52.0         1274.0          235.0
4      -122.25      37.85                52.0         1627.0          280.0

      population  households  median_income  median_house_value  \
0           322.0        126.0         8.3252         452600.0
1          2401.0       1138.0         8.3014        358500.0
2           496.0        177.0         7.2574        352100.0
3           558.0        219.0         5.6431        341300.0
4           565.0        259.0         3.8462        342200.0

      ocean_proximity_<1H OCEAN  ocean_proximity_INLAND  ocean_proximity_ISLAND  \
0                        0                0                0
1                        0                0                0
2                        0                0                0
3                        0                0                0
4                        0                0                0

      ocean_proximity_NEAR BAY  ocean_proximity_NEAR OCEAN
0                        1                0
1                        1                0
2                        1                0
3                        1                0
4                        1                0

```

```

In [63]: feat_removed = ['median_house_value']

```

```

# removed
#['longitude', 'latitude', 'housing_median_age', 'total_rooms', 'total_bedrooms', 'po
# 'median_house_value', 'ocean_proximity']

```

```

In [64]: y = df.median_house_value
X = df.drop(columns=feat_removed)
X.shape, y.shape

```

```

Out [64]: ((20640, 13), (20640,))

```

```

In [65]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

```

### 3.2 Metric RMSE root mean squared error

Root Mean Square Error (RMSE) is the standard deviation of the residuals (prediction errors). Residuals are a measure of how far from the regression line data points are; RMSE is a measure of how spread out these residuals are. In other words, it tells you how concentrated the data is around the line of best fit. Root mean square error is commonly used in climatology, forecasting, and regression analysis to verify experimental results.

```
In [66]: def calculate_rmse(model, model_name):
        model.fit(X_train, y_train)
        y_pred, y_pred_train = model.predict(X_test), model.predict(X_train)
        rmse_test, rmse_train = np.sqrt(mean_squared_error(y_test, y_pred)), np.sqrt(mean_squared_error(y_train, y_pred_train))
        print(model_name, f' RMSE on train: {rmse_train:.0f}, on test: {rmse_test:.0f}')
        return rmse_test
```

### 3.3 Linear Regression

```
In [67]: lr = LinearRegression()
        lr_err = calculate_rmse(lr, 'Linear Reg')
```

Linear Reg RMSE on train: 68533, on test: 69932

### 3.4 RANSAC Regressor

```
In [68]: ra = RANSACRegressor()
        ra_err = calculate_rmse(ra, 'RANSAC Reg')
```

RANSAC Reg RMSE on train: 78281, on test: 78795

### 3.5 Lasso

```
In [69]: la = Lasso()
        la_err = calculate_rmse(la, 'Lasso Reg')
```

Lasso Reg RMSE on train: 68533, on test: 69932

```
/home/sunflowa/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/coordinate_descent.py:176:
ConvergenceWarning
```

### 3.6 SGD Regressor

```
In [70]: sg = SGDRegressor()
        sg_err = calculate_rmse(sg, 'SGD Reg')
```

SGD Reg RMSE on train: 2939589866401599, on test: 2954302199978100

```
/home/sunflowa/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/stochastic_gradient_descent.py:176:
FutureWarning
```

### 3.7 Ridge

```
In [71]: ri = SGDRegressor()  
         ri_err = calculate_rmse(ri, 'Ridge')
```

Ridge RMSE on train: 24125952802617160, on test: 24256192059939448

```
/home/sunflowa/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/stochastic_gradient.  
FutureWarning)
```

### 3.8 AdaBoostRegressor

```
In [72]: ad = AdaBoostRegressor()  
         ad_err = calculate_rmse(ad, 'AdaBoostRegressor')
```

AdaBoostRegressor RMSE on train: 86734, on test: 86345

### 3.9 SVR


```
In [73]: sv = SVR()  
         sv_err = calculate_rmse(sv, 'SVR')
```

```
/home/sunflowa/anaconda3/lib/python3.7/site-packages/sklearn/svm/base.py:196: FutureWarning: T  
"avoid this warning.", FutureWarning)
```

SVR RMSE on train: 118660, on test: 117282

### 3.10 Results comparison

```
In [74]: df_score = pd.DataFrame({'Model': ['Linear Reg', 'RANSAC Reg', 'Lasso Reg', 'AdaBoost'  
                                           'RMSE': [lr_err, ra_err, la_err, ad_err, sv_err]})  
         ax = df_score.plot.barh(y='RMSE', x='Model')
```



Lasso and the Linear Reg are the winners ! Surprisingly the RSME is a little lower for the best models when we keep features such as lat/long and 'total\_bedrooms', 'population'.

In [ ]: